

# 01\_data\_acquisition

December 8, 2025

## 1 Data Acquisition and Optimized Sampling

This notebook handles loading the Criteo Display Advertising Challenge dataset, performing basic exploratory data analysis, and creating a **fatigue-optimized sample** for analysis.

### 1.1 Steps:

1. Load the Criteo dataset (from HuggingFace or local file)
2. Inspect data structure and perform basic EDA
3. Analyze multi-exposure users in the full dataset
4. Create fatigue-optimized sample (enriched with multi-exposure users)
5. Verify exposure distribution in the optimized sample
6. Save sample with metadata

### 1.2 Why Optimized Sampling?

The original 1% random sample has **only 2.25% multi-exposure users**, which severely limits fatigue analysis. This notebook creates a **fatigue-optimized sample** that:

- Identifies all multi-exposure users in the full dataset
- Samples 50% from multi-exposure users (instead of 2.25%)
- Samples 50% from single-exposure users (to maintain baseline)
- Preserves click distribution within each group

**Expected Improvement:** - Multi-exposure users: 2.25% → ~50% - Reliable exposure levels: 2 → 3-6+ - Statistical power: Low → High

```
[1]: import sys
import os
sys.path.append('..')

import pandas as pd
import numpy as np
import json
from src.data_loader import (
    load_criteo_data,
    inspect_data_structure,
    load_config
)
from src.smart_sampling import (
    identify_multi_exposure_users,
```

```

        create_fatigue_optimized_sample,
        save_optimized_sample
    )
from pathlib import Path

# Load configuration
config = load_config('../config/config.yaml')
print("Configuration loaded:")
print(f"  Random seed: {config['sampling']['random_seed']}")
```

Configuration loaded:

Random seed: 42

### 1.3 Step 1: Load Criteo Dataset

The dataset can be loaded from: - HuggingFace Hub (if available) - Local file path

**Note:** Update the data\_path below if you have the dataset locally.

```
[2]: # Load Criteo Attribution Dataset (TSV format)
# This dataset has: timestamp, uid, campaign, click, conversion, and contextual
# ↪features
data_path = "../data/raw/criteo_attribution_dataset.tsv.gz"

if os.path.exists(data_path):
    print(f"Loading Criteo Attribution Dataset from: {data_path}")
    print("This dataset contains:")
    print("  - timestamp: impression timestamp")
    print("  - uid: unique user identifier")
    print("  - campaign: campaign identifier")
    print("  - click: click label (0/1)")
    print("  - conversion: conversion label (0/1)")
    print("  - cat1-9: contextual features")
    print("  - And other attribution-related fields")
    df = load_criteo_data(data_path=data_path, use_huggingface=False)
else:
    print(f"Error: Dataset not found at {data_path}")
    print("Please ensure the file exists in the data/raw directory")
```

Loading Criteo Attribution Dataset from:

../data/raw/criteo\_attribution\_dataset.tsv.gz

This dataset contains:

- timestamp: impression timestamp
- uid: unique user identifier
- campaign: campaign identifier
- click: click label (0/1)
- conversion: conversion label (0/1)
- cat1-9: contextual features
- And other attribution-related fields

```

Loading data from ../data/raw/criteo_attribution_dataset.tsv.gz...
Detected TSV.GZ format, loading with tab separator...
Loaded 16,468,027 records
Columns: ['timestamp', 'uid', 'campaign', 'conversion', 'conversion_timestamp',
'conversion_id', 'attribution', 'click', 'click_pos', 'click_nb']...

```

## 1.4 Step 2: Inspect Data Structure

```
[3]: # Inspect data structure
if 'df' in locals():
    info = inspect_data_structure(df)
    print("Data Structure Information:")
    print(f"Shape: {info['shape']}")
    print(f"Memory usage: {info['memory_usage_mb']:.2f} MB")
    print(f"\nColumns ({len(info['columns'])}):")
    for col in info['columns'][:10]: # Show first 10
        print(f" - {col}: {info['dtypes'][col]}")
    if len(info['columns']) > 10:
        print(f" ... and {len(info['columns']) - 10} more columns")

    print(f"\nMissing values:")
    missing = {k: v for k, v in info['missing_values'].items() if v > 0}
    if missing:
        for col, count in list(missing.items())[:10]:
            print(f" {col}: {count}")
    else:
        print(" No missing values")

    # Display first few rows
    print("\nFirst few rows:")
    display(df.head())

```

Data Structure Information:

Shape: (16468027, 22)

Memory usage: 2764.10 MB

Columns (22):

- timestamp: int64
- uid: int64
- campaign: int64
- conversion: int64
- conversion\_timestamp: int64
- conversion\_id: int64
- attribution: int64
- click: int64
- click\_pos: int64
- click\_nb: int64
- ... and 12 more columns

```
Missing values:
```

```
    No missing values
```

```
First few rows:
```

```
      timestamp        uid  campaign  conversion  conversion_timestamp \
0            0  20073966  22589171          0             -1
1            2  24607497  884761          0             -1
2            2  28474333 18975823          0             -1
3            3  7306395  29427842          1            1449193
4            3  25357769 13365547          0             -1

  conversion_id  attribution  click  click_pos  click_nb ... \
0           -1           0     0       -1       -1   ...
1           -1           0     0       -1       -1   ...
2           -1           0     0       -1       -1   ...
3          3063962          0     1       0       7   ...
4           -1           0     0       -1       -1   ...

      time_since_last_click      cat1      cat2      cat3      cat4      cat5 \
0                  -1  5824233  9312274  3490278  29196072 11409686
1                 423858 30763035  9312274 14584482  29196072 11409686
2                  8879  138937  9312274 10769841  29196072 5824237
3                  -1  28928366 26597095 12435261 23549932 5824237
4                  -1  138937 26597094 31616034  29196072 11409684

      cat6      cat7      cat8      cat9
0  1973606  25162884  29196072  29196072
1  1973606  22644417  9312274 21091111
2  138937  1795451  29196072 15351056
3  1973606  9180723  29841067  29196072
4 26597096  4480345  29196072  29196072
```

```
[5 rows x 22 columns]
```

## 1.5 Step 3: Create Stratified Sample

```
[5]: # Analyze multi-exposure users in full dataset
print("=" * 60)
print("MULTI-EXPOSURE USER ANALYSIS (FULL DATASET)")
print("=" * 60)

if 'df' in locals():
    multi_exp_users, stats = identify_multi_exposure_users(
        df, user_col='uid', campaign_col='campaign', min_exposures=2
    )
```

```

print(f"\nTotal users: {stats['total_users']}")  

print(f"Multi-exposure users: {stats['multi_exposure_users']}")  

print(f"Percentage: {stats['pct_multi_exposure']:.2f}%")  

print(f"Max exposures: {stats['max_exposures_in_data']}")  
  

print(f"\nExposure distribution (max per user):")  

for exp, count in sorted(stats['exposure_distribution'].items())[:15]:  

    pct = count / stats['total_users'] * 100  

    print(f" {exp} exposures: {count}, users ({pct:.2f}%)")  

else:  

    print("Error: Dataset not loaded. Please run the data loading cell first.")  

=====  

MULTI-EXPOSURE USER ANALYSIS (FULL DATASET)  

=====
```

Total users: 6,142,256  
 Multi-exposure users: 2,545,854  
 Percentage: 41.45%  
 Max exposures: 376

Exposure distribution (max per user):  
 1 exposures: 3,596,402 users (58.55%)  
 2 exposures: 1,123,493 users (18.29%)  
 3 exposures: 510,431 users (8.31%)  
 4 exposures: 281,325 users (4.58%)  
 5 exposures: 172,172 users (2.80%)  
 6 exposures: 113,667 users (1.85%)  
 7 exposures: 78,840 users (1.28%)  
 8 exposures: 56,946 users (0.93%)  
 9 exposures: 41,578 users (0.68%)  
 10 exposures: 31,420 users (0.51%)

## 1.6 Step 4: Create Fatigue-Optimized Sample

Now we create a fatigue-optimized sample that enriches the dataset with multi-exposure users while maintaining a balanced representation.

```
[6]: # Create optimized sample  

print("=" * 60)  

print("CREATING FATIGUE-OPTIMIZED SAMPLE")  

print("=" * 60)  
  

if 'df' in locals():  

    sample, metadata = create_fatigue_optimized_sample(  

        df,  

        user_col='uid',  

        campaign_col='campaign',
```

```

        click_col='click',
        target_sample_size=500000,    # Target sample size
        multi_exp_ratio=0.5,         # 50% from multi-exposure users
        min_exposures=2,
        random_seed=config['sampling']['random_seed']
    )

    print(f"\n Sample created successfully!")
    print(f"    Sample size: {metadata['sample_size']:,}")
    print(f"    Multi-exposure users: {metadata['sample_multi_exp_pct']:.1f}%")
    print(f"    Click rate: {metadata['sample_click_rate']:.4f}")

else:
    print("Error: Dataset not loaded. Please run the data loading cell first.")

=====

```

CREATING FATIGUE-OPTIMIZED SAMPLE

Creating fatigue-optimized sample...

Target size: 500,000

Multi-exposure ratio: 50%

Multi-exposure users in full data: 2,545,854 (41.4%)

Multi-exposure records: 12,393,844

Single-exposure records: 4,074,183

Sampled 251,881 multi-exposure records

Sampled 248,119 single-exposure records

Final sample: 500,000 records

Multi-exposure users: 51,353 (17.3%)

Click rate: 0.3334

Sample created successfully!

Sample size: 500,000

Multi-exposure users: 17.3%

Click rate: 0.3334

## 1.7 Step 5: Verify Exposure Distribution in Optimized Sample

We verify that the optimized sample has sufficient data at multiple exposure levels for reliable fatigue analysis.

```
[8]: # Compute exposure counts in the optimized sample
if 'sample' in locals():
    sample_sorted = sample.sort_values(['uid', 'campaign', 'timestamp'])
    sample_sorted['exposure_count'] = (
        sample_sorted.groupby(['uid', 'campaign']).cumcount() + 1
    )
```

```

exp_dist = sample_sorted['exposure_count'].value_counts().sort_index()

print("EXPOSURE DISTRIBUTION IN OPTIMIZED SAMPLE:")
print("-" * 50)
print(f"{'Exposure':<15} {'Records':<15} {'% of Total':<15}")
print("-" * 50)

for exp in exp_dist.index[:15]:
    count = exp_dist[exp]
    pct = count / len(sample_sorted) * 100
    reliable = "Reliable" if count >= 100 else ""
    print(f"{exp:<15} {count:,} {pct:.2f}% {reliable}")

reliable_levels = sum(exp_dist >= 100)
print(f"\nReliable exposure levels (n>=100): {reliable_levels}")
print(f"Max exposure level: {exp_dist.index.max()}")
else:
    print("Error: Sample not created. Please run the sampling cell first.")

```

EXPOSURE DISTRIBUTION IN OPTIMIZED SAMPLE:

Exposure	Records	% of Total	
1	323,705	64.74%	Reliable
2	58,648	11.73%	Reliable
3	31,720	6.34%	Reliable
4	19,929	3.99%	Reliable
5	13,664	2.73%	Reliable
6	9,845	1.97%	Reliable
7	7,353	1.47%	Reliable
8	5,717	1.14%	Reliable
9	4,535	0.91%	Reliable
10	3,607	0.72%	Reliable
11	2,952	0.59%	Reliable
12	2,431	0.49%	Reliable
13	1,983	0.40%	Reliable
14	1,675	0.34%	Reliable
15	1,393	0.28%	Reliable

Reliable exposure levels (n>=100): 38

Max exposure level: 196

## 1.8 Step 6: Compare with Original Sample (if available)

If an original random sample exists, we compare it with the optimized sample to show the improvement.

```
[10]: # Compare with original sample if it exists
original_meta_path = "../data/samples/criteo_sample_01_metadata.json"

if os.path.exists(original_meta_path) and 'metadata' in locals():
    print("=" * 60)
    print("COMPARISON: ORIGINAL vs OPTIMIZED SAMPLE")
    print("=" * 60)

    with open(original_meta_path, 'r') as f:
        original_meta = json.load(f)

        print(f"\n{'Metric':<35} {'Original':<15} {'Optimized':<15}")
        print("-" * 65)
        print(f"{'Sample size':<35} {original_meta['sample_size']:,}")
        print(f"{'Multi-exp users (%)':<35} {original_meta['sample_multi_exp_pct']:.1f}%")
        print(f"{'Click rate':<35} {original_meta['click_rate']:.4f}")
        print(f"{'sample_click_rate':<35} {original_meta['sample_click_rate']:.4f}")

        improvement = metadata['sample_multi_exp_pct'] / original_meta['sample_multi_exp_pct']
        print(f"\nMulti-exposure users increased by {improvement:.1f}x!")

    else:
        print("Original sample metadata not found. Skipping comparison.")

=====
COMPARISON: ORIGINAL vs OPTIMIZED SAMPLE
=====
```

Metric	Original	Optimized
Sample size	164,681	500,000
Multi-exp users (%)	2.2%	17.3%
Click rate	0.3612	0.3334

Multi-exposure users increased by 7.7x!

## 1.9 Step 7: Save Optimized Sample

Save the fatigue-optimized sample and its metadata for use in subsequent analysis notebooks.

```
[11]: # Save the optimized sample
if 'sample' in locals() and 'metadata' in locals():
    sample_file, metadata_file = save_optimized_sample()
```

```

        sample,
        metadata,
        output_dir="../data/samples",
        sample_name="criteo_fatigue_optimized"
    )

    print(f"\n Optimized sample saved!")
    print(f"  Sample file: {sample_file}")
    print(f"  Metadata file: {metadata_file}")
    print(f"\n  To use this sample in other notebooks, load it with:")
    print(f"  df = pd.read_csv('../data/samples/criteo_fatigue_optimized.
        ↪csv')")
else:
    print("Error: Sample not created. Please run the sampling cell first.")

```

Sample saved to ../data/samples/criteo\_fatigue\_optimized.csv  
Metadata saved to ../data/samples/criteo\_fatigue\_optimized\_metadata.json

Optimized sample saved!  
Sample file: ../data/samples/criteo\_fatigue\_optimized.csv  
Metadata file: ../data/samples/criteo\_fatigue\_optimized\_metadata.json

To use this sample in other notebooks, load it with:  
df = pd.read\_csv('../data/samples/criteo\_fatigue\_optimized.csv')

[ ]: