

# 02\_exploratory\_analysis

December 8, 2025

## 1 Exploratory Data Analysis

This notebook performs comprehensive exploratory data analysis to understand:

- Data quality and basic statistics
- Temporal patterns in CTR
- Campaign-level characteristics
- Preliminary exposure analysis (prep for Research Question 1)

---

### 1.1 Prerequisites

**Input Data:** `../data/samples/criteo_fatigue_optimized.csv` (created by `01_data_acquisition.ipynb`)

**Output Data:** `../data/processed/data_with_exposures.csv` (input for `03_feature_engineering.ipynb`)

### 1.2 Data Pipeline Position

`01_data_acquisition.ipynb` → `[02_exploratory_analysis.ipynb]` → `03_feature_engineering.ipynb` → `04_model_training.ipynb`

```
[1]: import sys
sys.path.append('..')

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

# Set style
sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (12, 6)

# Load sample data
sample_path = "../data/samples/criteo_fatigue_optimized.csv"
df = pd.read_csv(sample_path)

print(f"Loaded {len(df)} records")
print(f"Columns: {len(df.columns)}")
df.head()
```

```
Loaded 500,000 records
```

```
Columns: 22
```

```
[1]:    timestamp      uid  campaign  conversion  conversion_timestamp  \
0          7  5588915  27491436          0                 -1
1         37  15687359  2576437          0                 -1
2         97  7327454  28137208          0                 -1
3        110  22075924  32368244          0                 -1
4        141  23042801  32368244          0                 -1

   conversion_id  attribution  click  click_pos  click_nb  ...  \
0           -1          0       0       -1       -1  ...
1           -1          0       1       -1       -1  ...
2           -1          0       0       -1       -1  ...
3           -1          0       1       -1       -1  ...
4           -1          0       1       -1       -1  ...

  time_since_last_click      cat1      cat2      cat3      cat4      cat5  \
0            -1  138937  9312274  4281154  29196072  28928366
1            -1  5824233  7477605  8562388  29196072  32440044
2            -1  30763035  26597095  29562678  29196072  32440047
3            138  25259032  7477605  8516481  17152695  11409684
4            -1  138937  7477605  19107905  29196072  11409684

      cat6      cat7      cat8      cat9
0  29196072  21857352  29196072  29196072
1  1973606   6954913  32440044  29196072
2  30763035  12831389  29196072  29196072
3  1973606   6266944  29841067  22419790
4  1973606   17502812  32440044  29196072
```

[5 rows x 22 columns]

### 1.3 Step 1: Basic Data Quality and Statistics

```
[2]: # Identify key columns based on Criteo Attribution Dataset structure
# From README.md: timestamp, uid, campaign, click, conversion, cat1-9, etc.

# Column mappings for Criteo Attribution Dataset:
click_col = 'click' # Primary target: click (0/1)
# Alternative: use 'conversion' for conversion analysis

user_col = 'uid' # User identifier (from dataset)
campaign_col = 'campaign' # Campaign identifier (from dataset)
time_col = 'timestamp' # Timestamp (from dataset, starting from 0)

# Basic statistics
```

```

print("==== Basic Statistics ===")
print(f"Total impressions: {len(df)}")
if click_col in df.columns:
    print(f"Overall CTR: {df[click_col].mean():.4f} ({df[click_col].sum():,} clicks)")
    print(f"Click distribution:")
    print(df[click_col].value_counts())

print(f"\nMemory usage: {df.memory_usage(deep=True).sum() / 1024**2:.2f} MB")
print(f"\nMissing values:")
missing = df.isnull().sum()
missing = missing[missing > 0]
if len(missing) > 0:
    print(missing)
else:
    print("No missing values")

```

```

==== Basic Statistics ===
Total impressions: 500,000
Overall CTR: 0.3334 (166,716 clicks)
Click distribution:
0      333284
1      166716
Name: click, dtype: int64

```

Memory usage: 83.92 MB

Missing values:  
No missing values

## 1.4 Step 2: Temporal Patterns

```

[3]: if time_col in df.columns and click_col in df.columns:
        # Parse timestamp (if not already datetime)
        print(f"Original timestamp dtype: {df[time_col].dtype}")
        print(f"Timestamp sample (first 5): {df[time_col].head()}")
        print(f"Timestamp range: {df[time_col].min()} to {df[time_col].max()}")

        if not pd.api.types.is_datetime64_any_dtype(df[time_col]):
            # If timestamp is numeric (seconds since start), convert to datetime
            if df[time_col].dtype in ['int64', 'float64', 'int32', 'float32']:
                # Check if values are very large (might be milliseconds or
                # microseconds)
                max_val = df[time_col].max()
                if max_val > 86400 * 365: # More than a year in seconds
                    # Might be milliseconds
                    if max_val > 86400 * 365 * 1000:

```

```

        # Might be microseconds
        df[time_col] = pd.to_datetime(df[time_col], unit='us')
    else:
        df[time_col] = pd.to_datetime(df[time_col], unit='ms')
    else:
        # Assume seconds
        reference_date = pd.Timestamp('2024-01-01')
        df[time_col] = pd.to_datetime(reference_date) + pd.
        to_timedelta(df[time_col], unit='s')
    else:
        df[time_col] = pd.to_datetime(df[time_col])

    print(f"\nConverted timestamp dtype: {df[time_col].dtype}")
    print(f"Converted timestamp range: {df[time_col].min()} to {df[time_col].
        max()}")


# Extract temporal features
df['hour'] = df[time_col].dt.hour
df['day_of_week'] = df[time_col].dt.dayofweek
df['day_name'] = df[time_col].dt.day_name()
df['month'] = df[time_col].dt.month

print(f"\nHour range: {df['hour'].min()} to {df['hour'].max()}")
print(f"Unique hours: {sorted(df['hour'].unique())[:10]}...") # Show first 10
print(f"Day names present: {df['day_name'].unique()}")


# CTR by hour
ctr_by_hour = df.groupby('hour')[click_col].agg(['mean', 'count'])
ctr_by_hour.columns = ['CTR', 'Impressions']
ctr_by_hour = ctr_by_hour.sort_index() # Sort by hour


# CTR by day of week
ctr_by_dow = df.groupby('day_name')[click_col].agg(['mean', 'count'])
ctr_by_dow.columns = ['CTR', 'Impressions'] # Rename columns first
# Reindex to ensure all days are present (fill missing with NaN)
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
    'Saturday', 'Sunday']
ctr_by_dow = ctr_by_dow.reindex(day_order)

print(f"\nCTR by hour summary:")
print(ctr_by_hour.head(10))
print(f"\nCTR by day of week summary:")
print(ctr_by_dow)


# Visualize
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

```

```

# Hour plot
axes[0].plot(ctr_by_hour.index, ctr_by_hour['CTR'], marker='o',  

             linewidth=2, markersize=6)
axes[0].set_xlabel('Hour of Day', fontsize=11)
axes[0].set_ylabel('CTR', fontsize=11)
axes[0].set_title('CTR by Hour of Day', fontsize=12, fontweight='bold')
axes[0].set_xlim(-0.5, 23.5) # Set proper x-axis range for hours
axes[0].set_xticks(range(0, 24, 2)) # Show every 2 hours
axes[0].grid(True, alpha=0.3)

# Day of week plot
ctr_by_dow_clean = ctr_by_dow.dropna()
if len(ctr_by_dow_clean) > 0:
    axes[1].bar(range(len(ctr_by_dow_clean)), ctr_by_dow_clean['CTR'],  

                alpha=0.7)
    axes[1].set_xticks(range(len(ctr_by_dow_clean)))
    axes[1].set_xticklabels(ctr_by_dow_clean.index, rotation=45, ha='right')
    axes[1].set_ylabel('CTR', fontsize=11)
    axes[1].set_title('CTR by Day of Week', fontsize=12, fontweight='bold')
else:
    axes[1].text(0.5, 0.5, 'No data available', ha='center', va='center',  

                transform=axes[1].transAxes)
    axes[1].set_title('CTR by Day of Week', fontsize=12, fontweight='bold')
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig('../results/figures/temporal_patterns.png', dpi=300,  

            bbox_inches='tight')
plt.show()

print("\nTemporal patterns saved to ../results/figures/temporal_patterns.  

      png")
else:
    print(f"Missing required columns. time_col='{time_col}',  

          click_col='{click_col}'")
    print(f"Available columns: {list(df.columns)}")

```

```

Original timestamp dtype: int64
Timestamp sample (first 5): 0      7
1     37
2     97
3    110
4    141
Name: timestamp, dtype: int64
Timestamp range: 7 to 2671193

```

```
Converted timestamp dtype: datetime64[ns]
Converted timestamp range: 2024-01-01 00:00:07 to 2024-01-31 21:59:53
```

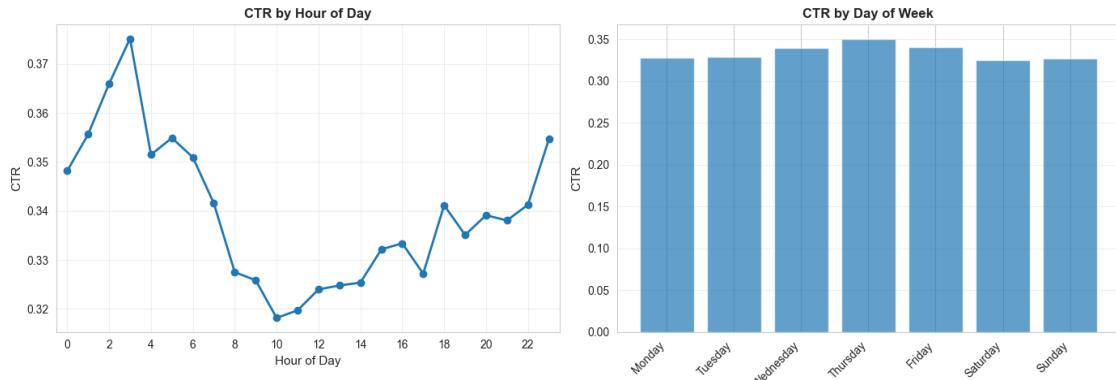
```
Hour range: 0 to 23
Unique hours: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]...
Day names present: ['Monday' 'Tuesday' 'Wednesday' 'Thursday' 'Friday'
'Saturday' 'Sunday']
```

CTR by hour summary:

hour	CTR	Impressions
0	0.348179	4064
1	0.355663	2896
2	0.365994	2776
3	0.375207	3622
4	0.351493	6094
5	0.354923	11101
6	0.350956	17042
7	0.341594	22758
8	0.327515	25953
9	0.325852	26644

CTR by day of week summary:

day_name	CTR	Impressions
Monday	0.327658	80819
Tuesday	0.329273	79806
Wednesday	0.339333	75112
Thursday	0.350126	56174
Friday	0.340403	68557
Saturday	0.324873	72533
Sunday	0.326871	66999



Temporal patterns saved to ../results/figures/temporal\_patterns.png

## 1.5 Step 3: Campaign-Level Analysis

```
[4]: if campaign_col in df.columns and click_col in df.columns:
    # Campaign statistics
    campaign_stats = df.groupby(campaign_col).agg({
        click_col: ['count', 'sum', 'mean']
    }).reset_index()
    campaign_stats.columns = [campaign_col, 'impressions', 'clicks', 'ctr']

    print(f"Total campaigns: {campaign_stats[campaign_col].nunique():,}")
    print(f"\nCampaign size distribution:")
    print(campaign_stats['impressions'].describe())

    print(f"\nCampaign CTR distribution:")
    print(campaign_stats['ctr'].describe())

    # Top campaigns by volume
    print("\nTop 10 campaigns by impressions:")
    print(campaign_stats.nlargest(10, 'impressions')[[campaign_col, 'impressions', 'ctr']])

    # Top campaigns by CTR
    print("\nTop 10 campaigns by CTR (min 100 impressions):")
    top_ctr = campaign_stats[campaign_stats['impressions'] >= 100].nlargest(10, 'ctr')
    print(top_ctr[[campaign_col, 'impressions', 'ctr']])

    # Visualize
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    axes[0].hist(campaign_stats['impressions'], bins=50, edgecolor='black', alpha=0.7)
    axes[0].set_xlabel('Impressions per Campaign')
    axes[0].set_ylabel('Number of Campaigns')
    axes[0].set_title('Campaign Size Distribution')
    axes[0].set_yscale('log')
    axes[0].grid(True, alpha=0.3)

    axes[1].hist(campaign_stats['ctr'], bins=50, edgecolor='black', alpha=0.7)
    axes[1].set_xlabel('CTR')
    axes[1].set_ylabel('Number of Campaigns')
    axes[1].set_title('Campaign CTR Distribution')
    axes[1].grid(True, alpha=0.3)

    plt.tight_layout()
```

```

plt.savefig('../results/figures/campaign_analysis.png', dpi=300,
bbox_inches='tight')
plt.show()

```

Total campaigns: 675

Campaign size distribution:

count	675.000000
mean	740.740741
std	1347.228097
min	1.000000
25%	149.500000
50%	322.000000
75%	753.500000
max	14012.000000

Name: impressions, dtype: float64

Campaign CTR distribution:

count	675.000000
mean	0.319524
std	0.098408
min	0.000000
25%	0.258431
50%	0.314775
75%	0.371690
max	0.691729

Name: ctr, dtype: float64

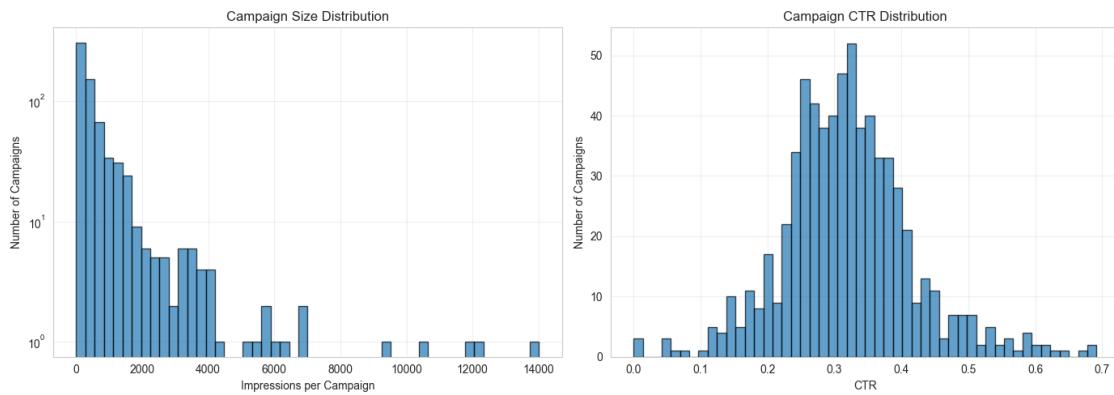
Top 10 campaigns by impressions:

	campaign	impressions	ctr
640	30801593	14012	0.382244
379	17686799	12054	0.294674
217	10341182	11849	0.409908
329	15398570	10445	0.332791
101	5061834	9519	0.267150
598	29427842	6942	0.306540
326	15184511	6822	0.324978
399	18975823	6349	0.439597
575	28351001	5983	0.588166
10	497593	5831	0.396501

Top 10 campaigns by CTR (min 100 impressions):

	campaign	impressions	ctr
634	30534043	266	0.691729
279	13365547	1106	0.673599
243	11959615	1151	0.637706
192	9095251	982	0.627291
627	30405203	274	0.620438

265	12843316	272	0.610294
53	2946551	1448	0.598066
209	9810200	806	0.591811
497	24982961	176	0.590909
575	28351001	5983	0.588166



## 1.6 Step 4: Preliminary Exposure Analysis (Research Question 1 Prep)

This analysis provides an initial look at how CTR changes with exposure count.

```
[5]: if all(col in df.columns for col in [user_col, campaign_col, time_col, click_col]):
    # Sort by timestamp
    df_sorted = df.sort_values([user_col, campaign_col, time_col])

    # Compute exposure count for each user-campaign pair
    df_sorted['exposure_count'] = (
        df_sorted.groupby([user_col, campaign_col])
        .cumcount() + 1
    )

    print(f"Exposure count statistics:")
    print(f"  Max exposure: {df_sorted['exposure_count'].max()}")
    print(f"  Mean exposure: {df_sorted['exposure_count'].mean():.2f}")
    print(f"  Users with >1 exposure: {(df_sorted['exposure_count'] > 1).sum():,}")
    print(f"  User-campaign pairs: {df_sorted.groupby([user_col, campaign_col]).ngroups:,}")

    # Compute CTR by exposure count
    ctr_by_exposure = df_sorted.groupby('exposure_count')[click_col].agg([
        ('ctr', 'mean'),
        ('impressions', 'count'),
```

```

        ('clicks', 'sum')
    ]).reset_index()

    # Calculate confidence intervals (95% CI using normal approximation)
    n = ctr_by_exposure['impressions']
    p = ctr_by_exposure['ctr']
    se = np.sqrt(p * (1 - p) / n)
    z = 1.96 # 95% confidence
    ctr_by_exposure['ctr_lower'] = np.maximum(0, p - z * se)
    ctr_by_exposure['ctr_upper'] = np.minimum(1, p + z * se)

    # Filter to exposure counts with sufficient samples (lower threshold to see more data)
    min_samples = 50 # Reduced from 100 to see more exposure levels
    ctr_by_exposure_filtered = ctr_by_exposure[ctr_by_exposure['impressions'] >= min_samples].copy()

    print(f"\nCTR by Exposure Count (min {min_samples} samples per exposure level):")
    print(ctr_by_exposure_filtered.head(30))
    print(f"\nTotal exposure levels with sufficient data: {len(ctr_by_exposure_filtered)}")

    # Check if CTR is increasing or decreasing
    if len(ctr_by_exposure_filtered) >= 2:
        first_ctr = ctr_by_exposure_filtered.iloc[0]['ctr']
        last_ctr = ctr_by_exposure_filtered.iloc[-1]['ctr']
        change = last_ctr - first_ctr
        change_pct = (change / first_ctr * 100) if first_ctr > 0 else 0
        print(f"\nCTR change from exposure 1 to exposure {ctr_by_exposure_filtered['exposure_count'].max()}:")
        print(f" Exposure 1 CTR: {first_ctr:.4f}")
        print(f" Exposure {ctr_by_exposure_filtered['exposure_count'].max()} CTR: {last_ctr:.4f}")
        print(f" Change: {change:+.4f} ({change_pct:+.2f}%)")
        if change > 0:
            print(" Note: CTR is INCREASING with exposure (opposite of fatigue)")
            print(" This could indicate selection bias or that engaged users see more ads")
        else:
            print(" Note: CTR is DECREASING with exposure (consistent with fatigue)")

    # Visualize preliminary decay curve with confidence intervals
    plt.figure(figsize=(12, 6))

```

```

# Plot with confidence intervals
plt.plot(ctr_by_exposure_filtered['exposure_count'], □
         ↪ctr_by_exposure_filtered['ctr'],
         'o-', linewidth=2, markersize=6, label='Mean CTR', color='#2E86AB')
plt.fill_between(
    ctr_by_exposure_filtered['exposure_count'],
    ctr_by_exposure_filtered['ctr_lower'],
    ctr_by_exposure_filtered['ctr_upper'],
    alpha=0.3, color='#2E86AB', label='95% Confidence Interval'
)

plt.xlabel('Exposure Count', fontsize=12)
plt.ylabel('Click-Through Rate (CTR)', fontsize=12)
plt.title('Preliminary CTR Trend with Exposure Count', fontsize=14, □
         ↪fontweight='bold')
plt.legend()
plt.grid(True, alpha=0.3)

# Add annotation if we see increasing trend
if len(ctr_by_exposure_filtered) >= 2:
    first_ctr = ctr_by_exposure_filtered.iloc[0]['ctr']
    last_ctr = ctr_by_exposure_filtered.iloc[-1]['ctr']
    if last_ctr > first_ctr:
        max_exp = ctr_by_exposure_filtered['exposure_count'].max()
        plt.text(0.02, 0.98,
                 f'Note: CTR increases with exposure\n(may indicate□
         ↪selection bias)',
                 transform=plt.gca().transAxes,
                 verticalalignment='top',
                 bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5),
                 fontsize=9)

plt.tight_layout()
plt.savefig('../results/figures/preliminary_decay.png', dpi=300, □
         ↪bbox_inches='tight')
plt.show()

print("\nPreliminary decay curve saved to ../results/figures/□
         ↪preliminary_decay.png")

# Save processed data for next steps
df_sorted.to_csv('../data/processed/data_with_exposures.csv', index=False)
print("\nData with exposure counts saved to ../data/processed/□
         ↪data_with_exposures.csv")
else:

```

```

print("Required columns not found. Please update column names in earlier cells.")
print(f"Looking for: {[user_col, campaign_col, time_col, click_col]}")
print(f"Available columns: {list(df.columns)}")

```

Exposure count statistics:

Max exposure: 196  
 Mean exposure: 2.76  
 Users with >1 exposure: 176,295  
 User-campaign pairs: 323,705

CTR by Exposure Count (min 50 samples per exposure level):

	exposure_count	ctr	impressions	clicks	ctr_lower	ctr_upper
0	1	0.304296	323705	98502	0.302711	0.305881
1	2	0.340523	58648	19971	0.336688	0.344358
2	3	0.353909	31720	11226	0.348647	0.359172
3	4	0.364695	19929	7268	0.358012	0.371378
4	5	0.385026	13664	5261	0.376867	0.393185
5	6	0.395429	9845	3893	0.385771	0.405088
6	7	0.405821	7353	2984	0.394597	0.417045
7	8	0.421375	5717	2409	0.408575	0.434175
8	9	0.428004	4535	1941	0.413604	0.442405
9	10	0.433324	3607	1563	0.417152	0.449496
10	11	0.460366	2952	1359	0.442385	0.478346
11	12	0.466475	2431	1134	0.446643	0.486306
12	13	0.476046	1983	944	0.454064	0.498028
13	14	0.488358	1675	818	0.464419	0.512297
14	15	0.515434	1393	718	0.489190	0.541679
15	16	0.493220	1180	582	0.464694	0.521747
16	17	0.510574	993	507	0.479482	0.541666
17	18	0.539863	878	474	0.506895	0.572831
18	19	0.560425	753	422	0.524974	0.595876
19	20	0.546708	653	357	0.508525	0.584890
20	21	0.573451	565	324	0.532670	0.614233
21	22	0.626804	485	304	0.583759	0.669849
22	23	0.578199	422	244	0.531080	0.625318
23	24	0.585366	369	216	0.535098	0.635634
24	25	0.623100	329	205	0.570734	0.675466
25	26	0.608997	289	176	0.552736	0.665257
26	27	0.653992	263	172	0.596500	0.711484
27	28	0.629167	240	151	0.568055	0.690278
28	29	0.609091	220	134	0.544611	0.673571
29	30	0.666667	198	132	0.601004	0.732329

Total exposure levels with sufficient data: 47

CTR change from exposure 1 to exposure 47:

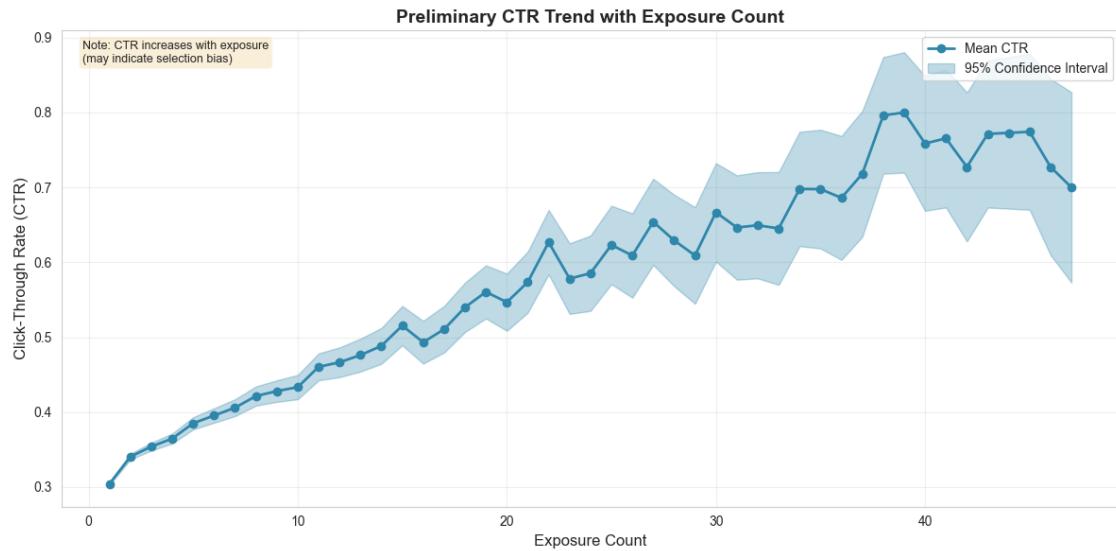
Exposure 1 CTR: 0.3043

Exposure 47 CTR: 0.7000

Change: +0.3957 (+130.04%)

Note: CTR is INCREASING with exposure (opposite of fatigue)

This could indicate selection bias or that engaged users see more ads



Preliminary decay curve saved to ../results/figures/preliminary\_decay.png

Data with exposure counts saved to ../data/processed/data\_with\_exposures.csv