

# 04\_rq1\_ctr\_fatigue\_analysis

December 8, 2025

## 1 Research Question 1: When Does Ad Fatigue Set In?

**Question:** How many impressions does it take before an ad's CTR meaningfully declines?

---

### 1.1 Prerequisites

**Input Data:** `../data/processed/data_with_all_features.csv` (created by `03_feature_engineering.ipynb`)  
**Output Data:** - `../results/figures/rq1_*.png` (visualizations) - `../results/tables/rq1_*.csv` (statistical results)

### 1.2 Data Pipeline Position

`01_data_acquisition` → `02_exploratory_analysis` → `03_feature_engineering` → [04\_rq1\_ctr\_fatigue\_anal

### 1.3 Source Module Dependencies

This notebook primarily uses **custom analysis functions defined inline** rather than src modules. This allows for transparent, reproducible analysis where all logic is visible in the notebook.

**Key functions defined in this notebook:** - `compute_cohort_ctr_by_exposure()`: Computes CTR at each exposure level for a fixed cohort - `test_ctr_difference()`: Statistical test for CTR decline between exposure levels - `fit_decay_curve()`: Fits exponential/power-law decay models

---

### 1.4 The Survivorship Bias Problem

**Why does CTR appear to INCREASE with exposure in naive analysis?**

When we simply group by exposure count and compute mean CTR, we're comparing different populations:  
- At exposure 1: ALL users (everyone starts with a first impression)  
- At exposure 10: ONLY users who stayed long enough to see 10 impressions

Users who see many ads are **inherently more engaged** (that's why they keep seeing ads!), so their CTR is higher.

## 1.5 Our Approach: Time-Bounded Within-User Analysis

To correctly measure fatigue, we:

1. **Use sequential exposure\_count** - the Nth impression for each user-campaign pair (computed chronologically)
2. **Compare within users** - same users at exposure 1 vs exposure N (controls for user-level confounders)
3. **Cohort analysis** - follow users who reached at least N exposures to avoid survivorship bias
4. **Time-bounded metrics** - compute CTR at each exposure level as a snapshot, not cumulative

```
[19]: import sys
sys.path.append('..')

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from typing import Dict, Tuple, Optional
import warnings
warnings.filterwarnings('ignore')

# Set plot style
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['font.size'] = 11

# Load data
df = pd.read_csv('../data/processed/data_with_all_features.csv')
print(f"Loaded {len(df)} impression records")
print(f"Unique users: {df['uid'].nunique()}")
print(f"Unique campaigns: {df['campaign'].nunique()}")
```

Loaded 500,000 impression records

Unique users: 297,407

Unique campaigns: 675

## 1.6 Step 1: Understanding the Exposure Count Column

The **exposure\_count** column represents the **Nth impression** for each user-campaign pair, computed chronologically. -  $\text{exposure\_count} = 1 \rightarrow$  This is the first time this user has seen this campaign -  $\text{exposure\_count} = 5 \rightarrow$  This is the 5th time this user has seen this campaign

This is the correct metric for fatigue analysis (not **exposure\_count\_168h** which is time-windowed and biased toward high-frequency users).

```
[21]: # Verify exposure_count is correct (Nth sequential impression)
print("=" * 70)
```

```

print("EXPOSURE COUNT DISTRIBUTION")
print("==" * 70)
print(f"\nExposure count statistics:")
print(df['exposure_count'].describe())

# Distribution of max exposures per user-campaign pair
user_campaign_max_exp = df.groupby(['uid', 'campaign'])['exposure_count'].max()
print(f"\nMax exposure per user-campaign pair:")
print(user_campaign_max_exp.describe())

# Percentage of impressions at each exposure level
exp_dist = df['exposure_count'].value_counts().sort_index().head(20)
exp_pct = (exp_dist / len(df) * 100).round(2)
print(f"\n% of impressions at each exposure level:")
for exp, pct in exp_pct.items():
    print(f"  Exposure {exp}: {pct}%")

```

=====

EXPOSURE COUNT DISTRIBUTION

=====

Exposure count statistics:

| count | 500000.00000 |
|-------|--------------|
| mean  | 2.762458     |
| std   | 5.942473     |
| min   | 1.000000     |
| 25%   | 1.000000     |
| 50%   | 1.000000     |
| 75%   | 2.000000     |
| max   | 196.000000   |

Name: exposure\_count, dtype: float64

Max exposure per user-campaign pair:

| count | 323705.00000 |
|-------|--------------|
| mean  | 1.544616     |
| std   | 2.145561     |
| min   | 1.000000     |
| 25%   | 1.000000     |
| 50%   | 1.000000     |
| 75%   | 1.000000     |
| max   | 196.000000   |

Name: exposure\_count, dtype: float64

% of impressions at each exposure level:

- Exposure 1: 64.74%
- Exposure 2: 11.73%
- Exposure 3: 6.34%
- Exposure 4: 3.99%

```

Exposure 5: 2.73%
Exposure 6: 1.97%
Exposure 7: 1.47%
Exposure 8: 1.14%
Exposure 9: 0.91%
Exposure 10: 0.72%
Exposure 11: 0.59%
Exposure 12: 0.49%
Exposure 13: 0.4%
Exposure 14: 0.34%
Exposure 15: 0.28%
Exposure 16: 0.24%
Exposure 17: 0.2%
Exposure 18: 0.18%
Exposure 19: 0.15%
Exposure 20: 0.13%

```

## 1.7 Step 2: Diagnose Survivorship Bias

First, let's demonstrate why naive CTR-by-exposure analysis is misleading.

```
[23]: print("=" * 70)
print("SURVIVORSHIP BIAS DIAGNOSIS")
print("=" * 70)

# Naive analysis: CTR by exposure count (THIS IS BIASED)
naive_ctr = df.groupby('exposure_count')['click'].agg(['mean', 'count']).  

    ↪reset_index()  

naive_ctr.columns = ['exposure_count', 'ctr', 'n_impressions']
naive_ctr = naive_ctr[naive_ctr['n_impressions'] >= 100] # Min samples

print("\nNAIVE CTR BY EXPOSURE (BIASED - DO NOT USE):")
print(f"{'Exp':<6} {'CTR':<10} {'N Impressions':<15}")
print("-" * 35)
for _, row in naive_ctr.head(15).iterrows():
    print(f"{int(row['exposure_count']):<6} {row['ctr']:.4f}      □  

    ↪{int(row['n_impressions']):,}"))

# Check trend
first_ctr = naive_ctr[naive_ctr['exposure_count'] == 1]['ctr'].values[0]
exp_5_ctr = naive_ctr[naive_ctr['exposure_count'] == 5]['ctr'].values[0]
exp_10_ctr = naive_ctr[naive_ctr['exposure_count'] == 10]['ctr'].values[0] if  

    ↪10 in naive_ctr['exposure_count'].values else None

print(f"\n  MISLEADING TREND (due to survivorship bias):")
print(f"  CTR at exposure 1: {first_ctr:.4f}")
print(f"  CTR at exposure 5: {exp_5_ctr:.4f} ({(exp_5_ctr/first_ctr - 1)*100:+.  

    ↪1f}% vs exp 1)"
```

```

if exp_10_ctr:
    print(f"    CTR at exposure 10: {exp_10_ctr:.4f} ({(exp_10_ctr/first_ctr - 1)*100:+.1f}% vs exp 1)")
print(f"\n    This makes it look like CTR INCREASES with exposure - clearly wrong!")

```

=====

## SURVIVORSHIP BIAS DIAGNOSIS

=====

NAIVE CTR BY EXPOSURE (BIASED - DO NOT USE):

| Exp | CTR    | N Impressions |
|-----|--------|---------------|
| 1   | 0.3043 | 323,705       |
| 2   | 0.3405 | 58,648        |
| 3   | 0.3539 | 31,720        |
| 4   | 0.3647 | 19,929        |
| 5   | 0.3850 | 13,664        |
| 6   | 0.3954 | 9,845         |
| 7   | 0.4058 | 7,353         |
| 8   | 0.4214 | 5,717         |
| 9   | 0.4280 | 4,535         |
| 10  | 0.4333 | 3,607         |
| 11  | 0.4604 | 2,952         |
| 12  | 0.4665 | 2,431         |
| 13  | 0.4760 | 1,983         |
| 14  | 0.4884 | 1,675         |
| 15  | 0.5154 | 1,393         |

MISLEADING TREND (due to survivorship bias):

```

CTR at exposure 1: 0.3043
CTR at exposure 5: 0.3850 (+26.5% vs exp 1)
CTR at exposure 10: 0.4333 (+42.4% vs exp 1)

```

This makes it look like CTR INCREASES with exposure - clearly wrong!

```

[25]: # Demonstrate the bias: Compare first-exposure CTR for different user types
print("\n" + "=" * 70)
print("WHY THE BIAS OCCURS")
print("=" * 70)

# Get max exposure per user-campaign pair
user_campaign_info = df.groupby(['uid', 'campaign']).agg({
    'exposure_count': 'max',
    'click': 'sum'
}).reset_index()
user_campaign_info.columns = ['uid', 'campaign', 'max_exposure', 'total_clicks']

```

```

# Get first exposure records
first_exposure_df = df[df['exposure_count'] == 1].merge(
    user_campaign_info[['uid', 'campaign', 'max_exposure']],
    on=['uid', 'campaign']
)

# Compare first-exposure CTR by eventual max exposure
print("\nFirst-exposure CTR by user's eventual total exposure:")
print("(If no bias, these should be similar)")
print(f"\n{'User Type':<30} {'First Exp CTR':<15} {'N Users':<15}")
print("-" * 60)

exposure_groups = [
    (1, 1, "1 impression only"),
    (2, 3, "2-3 impressions"),
    (4, 10, "4-10 impressions"),
    (11, 50, "11-50 impressions"),
    (51, 1000, "50+ impressions")
]
]

first_exp_ctrs = []
for min_exp, max_exp, label in exposure_groups:
    mask = (first_exposure_df['max_exposure'] >= min_exp) &
    ↪(first_exposure_df['max_exposure'] <= max_exp)
    group = first_exposure_df[mask]
    if len(group) >= 50:
        ctr = group['click'].mean()
        first_exp_ctrs.append((label, ctr, len(group)))
        print(f"{label:<30} {ctr:.4f} {len(group):,}")

print(f"\nINSIGHT:")
if len(first_exp_ctrs) >= 2:
    low_exp_ctr = first_exp_ctrs[0][1]
    high_exp_ctr = first_exp_ctrs[-1][1]
    if high_exp_ctr > low_exp_ctr:
        print(f"  Users who will see 50+ impressions have {high_exp_ctr/
    ↪low_exp_ctr:.2f}x higher CTR")
        print(f"  at their FIRST exposure ({high_exp_ctr:.4f} vs {low_exp_ctr:.
    ↪4f}).")
        print(f"  This confirms SELECTION BIAS: high-exposure users are
    ↪inherently more engaged.")

```

=====

WHY THE BIAS OCCURS

=====

First-exposure CTR by user's eventual total exposure:  
 (If no bias, these should be similar)

| User Type         | First Exp CTR | N Users |
|-------------------|---------------|---------|
| 1 impression only | 0.2732        | 265,057 |
| 2-3 impressions   | 0.4061        | 38,719  |
| 4-10 impressions  | 0.5125        | 16,977  |
| 11-50 impressions | 0.5598        | 2,912   |

INSIGHT:

Users who will see 50+ impressions have 2.05x higher CTR  
 at their FIRST exposure (0.5598 vs 0.2732).

This confirms SELECTION BIAS: high-exposure users are inherently more engaged.

## 1.8 Step 3: Correct Analysis - Within-User CTR Change

To correctly measure fatigue, we compare the **same users** at different exposure levels.

For each user-campaign pair with multiple exposures:  
 - CTR at exposure 1 = click rate at their first impression  
 - CTR at exposure N = click rate at their Nth impression

This controls for user-level confounders because we're comparing the same users to themselves.

```
[26]: print("=" * 70)
print("CORRECT ANALYSIS: WITHIN-USER CTR BY EXPOSURE")
print("=" * 70)

def compute_cohort_ctr_by_exposure(
    df: pd.DataFrame,
    min_exposures: int = 5,
    max_exposure_to_analyze: int = 20
) -> Tuple[pd.DataFrame, int]:
    """
    Compute CTR at each exposure level using cohort analysis.

    Only includes user-campaign pairs that reached at least min_exposures,
    ensuring we compare the same cohort across exposure levels.

    This is a time-bounded analysis: CTR at exposure N is the click rate
    of the Nth impression for each user-campaign pair.
    """
    # Get user-campaign pairs with sufficient exposures
    user_campaign_max = df.groupby(['uid', 'campaign'])['exposure_count'].max()
    cohort_pairs = user_campaign_max[user_campaign_max >= min_exposures].
    ↪reset_index()
    cohort_pairs = cohort_pairs[['uid', 'campaign']]
```

```

# Filter to cohort
cohort_df = df.merge(cohort_pairs, on=['uid', 'campaign'], how='inner')

# Compute CTR at each exposure level (time-bounded: CTR of Nth impression)
results = []
for exp in range(1, max_exposure_to_analyze + 1):
    exp_data = cohort_df[cohort_df['exposure_count'] == exp]
    if len(exp_data) >= 50:
        ctr = exp_data['click'].mean()
        n = len(exp_data)
        # Confidence interval
        se = np.sqrt(ctr * (1 - ctr) / n)
        ci_lower = max(0, ctr - 1.96 * se)
        ci_upper = min(1, ctr + 1.96 * se)
        results.append({
            'exposure': exp,
            'ctr': ctr,
            'n_impressions': n,
            'ci_lower': ci_lower,
            'ci_upper': ci_upper
        })
    else:
        print(f"Exposure {exp} has less than 50 impressions, skipping")

return pd.DataFrame(results), len(cohort_pairs)

# Analyze cohort with at least 5 exposures
cohort_ctr, cohort_size = compute_cohort_ctr_by_exposure(df, min_exposures=5, max_exposure_to_analyze=20)

print(f"\nCohort Analysis: Users with at least 5 exposures")
print(f"Cohort size: {cohort_size:,} user-campaign pairs")
print(f"\n{'Exposure':<10} {'CTR':<10} {'95% CI':<22} {'N Impressions':<15}")
print("-" * 60)
for _, row in cohort_ctr.iterrows():
    ci = f"[{row['ci_lower']:.4f}, {row['ci_upper']:.4f}]"
    print(f"{int(row['exposure']):<10} {row['ctr']:.4f}      {ci:<22} {int(row['n_impressions']):<15}")

=====
CORRECT ANALYSIS: WITHIN-USER CTR BY EXPOSURE
=====
```

Cohort Analysis: Users with at least 5 exposures  
Cohort size: 13,664 user-campaign pairs

| Exposure | CTR    | 95% CI           | N Impressions |
|----------|--------|------------------|---------------|
| 1        | 0.5285 | [0.5201, 0.5368] | 13,664        |

|    |        |                  |        |
|----|--------|------------------|--------|
| 2  | 0.4475 | [0.4391, 0.4558] | 13,664 |
| 3  | 0.4158 | [0.4075, 0.4240] | 13,664 |
| 4  | 0.3918 | [0.3836, 0.4000] | 13,664 |
| 5  | 0.3850 | [0.3769, 0.3932] | 13,664 |
| 6  | 0.3954 | [0.3858, 0.4051] | 9,845  |
| 7  | 0.4058 | [0.3946, 0.4170] | 7,353  |
| 8  | 0.4214 | [0.4086, 0.4342] | 5,717  |
| 9  | 0.4280 | [0.4136, 0.4424] | 4,535  |
| 10 | 0.4333 | [0.4172, 0.4495] | 3,607  |
| 11 | 0.4604 | [0.4424, 0.4783] | 2,952  |
| 12 | 0.4665 | [0.4466, 0.4863] | 2,431  |
| 13 | 0.4760 | [0.4541, 0.4980] | 1,983  |
| 14 | 0.4884 | [0.4644, 0.5123] | 1,675  |
| 15 | 0.5154 | [0.4892, 0.5417] | 1,393  |
| 16 | 0.4932 | [0.4647, 0.5217] | 1,180  |
| 17 | 0.5106 | [0.4795, 0.5417] | 993    |
| 18 | 0.5399 | [0.5069, 0.5728] | 878    |
| 19 | 0.5604 | [0.5250, 0.5959] | 753    |
| 20 | 0.5467 | [0.5085, 0.5849] | 653    |

```
[27]: # Calculate decline from exposure 1
if len(cohort_ctr) >= 3:
    baseline_ctr = cohort_ctr[cohort_ctr['exposure'] == 1]['ctr'].values[0]

    print("\n" + "=" * 70)
    print("CTR DECLINE FROM BASELINE (Exposure 1)")
    print("=" * 70)
    print(f"\nBaseline CTR (Exposure 1): {baseline_ctr:.4f}")
    print(f"\n{'Exposure':<10} {'CTR':<10} {'Absolute Change':<18} {'Relative Change':<18}")
    print("-" * 60)

    for _, row in cohort_ctr.iterrows():
        exp = int(row['exposure'])
        ctr = row['ctr']
        abs_change = ctr - baseline_ctr
        rel_change = (ctr - baseline_ctr) / baseline_ctr * 100 if baseline_ctr > 0 else 0
        print(f"{exp:<10} {ctr:.4f}      {abs_change:+.4f}          {rel_change:+.1f}%")

```

```
=====
CTR DECLINE FROM BASELINE (Exposure 1)
=====

Baseline CTR (Exposure 1): 0.5285
```

| Exposure | CTR    | Absolute Change | Relative Change |
|----------|--------|-----------------|-----------------|
| <hr/>    |        |                 |                 |
| 1        | 0.5285 | +0.0000         | +0.0%           |
| 2        | 0.4475 | -0.0810         | -15.3%          |
| 3        | 0.4158 | -0.1127         | -21.3%          |
| 4        | 0.3918 | -0.1366         | -25.9%          |
| 5        | 0.3850 | -0.1434         | -27.1%          |
| 6        | 0.3954 | -0.1330         | -25.2%          |
| 7        | 0.4058 | -0.1226         | -23.2%          |
| 8        | 0.4214 | -0.1071         | -20.3%          |
| 9        | 0.4280 | -0.1005         | -19.0%          |
| 10       | 0.4333 | -0.0951         | -18.0%          |
| 11       | 0.4604 | -0.0681         | -12.9%          |
| 12       | 0.4665 | -0.0620         | -11.7%          |
| 13       | 0.4760 | -0.0524         | -9.9%           |
| 14       | 0.4884 | -0.0401         | -7.6%           |
| 15       | 0.5154 | -0.0130         | -2.5%           |
| 16       | 0.4932 | -0.0352         | -6.7%           |
| 17       | 0.5106 | -0.0179         | -3.4%           |
| 18       | 0.5399 | +0.0114         | +2.2%           |
| 19       | 0.5604 | +0.0320         | +6.0%           |
| 20       | 0.5467 | +0.0182         | +3.5%           |

```
[28]: # Visualization: CTR by Exposure (Cohort Analysis)
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Left: CTR by exposure with confidence intervals
ax1 = axes[0]
ax1.plot(cohort_ctr['exposure'], cohort_ctr['ctr'], 'o-', color='#2E86AB', ▾
         linewidth=2, markersize=8)
ax1.fill_between(cohort_ctr['exposure'], cohort_ctr['ci_lower'], ▾
                 cohort_ctr['ci_upper'], ▾
                 alpha=0.3, color='#2E86AB')
ax1.axhline(y=baseline_ctr, color='#E94F37', linestyle='--', alpha=0.7, ▾
            label=f'Baseline CTR = {baseline_ctr:.4f}')
ax1.set_xlabel('Exposure Count (Nth Impression)', fontsize=12)
ax1.set_ylabel('Click-Through Rate (CTR)', fontsize=12)
ax1.set_title('CTR by Sequential Exposure\n(Same Cohort Tracked Over Time)', ▾
              fontsize=13, fontweight='bold')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Right: Relative decline
ax2 = axes[1]
rel_changes = [(row['ctr'] - baseline_ctr) / baseline_ctr * 100 for _, row in ▾
               cohort_ctr.iterrows()]
colors = ['#28A745' if c >= 0 else '#DC3545' for c in rel_changes]
```

```

ax2.bar(cohort_ctr['exposure'], rel_changes, color=colors, alpha=0.7,
       edgecolor='black')
ax2.axhline(y=0, color='black', linewidth=1)
ax2.axhline(y=-10, color='orange', linestyle='--', alpha=0.7, label='-10% threshold')
ax2.axhline(y=-20, color='red', linestyle='--', alpha=0.7, label='-20% threshold')
ax2.set_xlabel('Exposure Count (Nth Impression)', fontsize=12)
ax2.set_ylabel('% Change from Baseline', fontsize=12)
ax2.set_title('Relative CTR Change from First Exposure', fontsize=13,
              fontweight='bold')
ax2.legend()
ax2.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig('../results/figures/rq1_cohort_ctr_analysis.png', dpi=300,
            bbox_inches='tight')
plt.show()
print("\nFigure saved to ../results/figures/rq1_cohort_ctr_analysis.png")

```

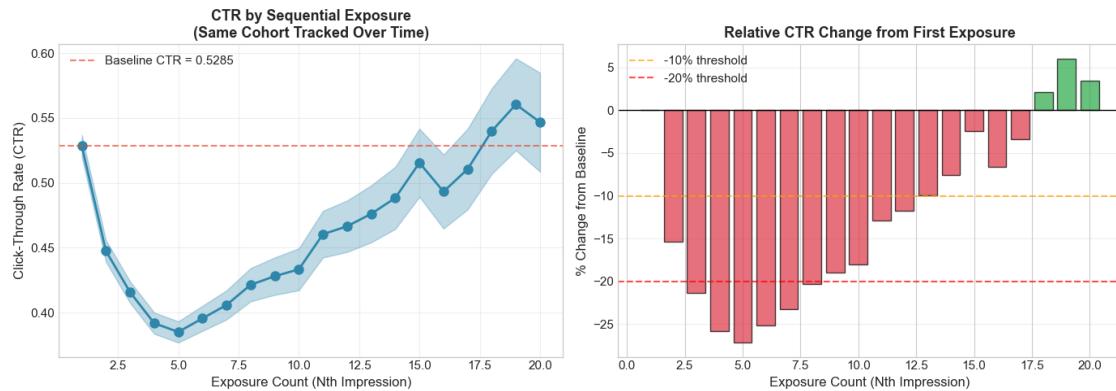


Figure saved to ../results/figures/rq1\_cohort\_ctr\_analysis.png

## 1.9 Step 4: Statistical Testing for CTR Decline

We test if the CTR decline is statistically significant using paired comparisons.

```
[29]: print("=" * 70)
print("STATISTICAL TESTING: IS CTR DECLINE SIGNIFICANT?")
print("=" * 70)

def test_ctr_decline_paired(
    df: pd.DataFrame,

```

```

        exposure_1: int,
        exposure_2: int,
        min_cohort_size: int = 5
    ) -> Dict:
        """
        Test CTR decline using paired comparison (same user-campaign pairs).
        This is a time-bounded test comparing the Nth impression CTR.
        """

        # Get pairs that have both exposures
        user_campaign_max = df.groupby(['uid', 'campaign'])['exposure_count'].max()
        cohort_pairs = user_campaign_max[user_campaign_max >= max(exposure_1, exposure_2)].reset_index()
        cohort_pairs = cohort_pairs[['uid', 'campaign']]

        if len(cohort_pairs) < min_cohort_size:
            return {'error': 'Insufficient cohort size'}

        # Get clicks at each exposure level for the cohort
        cohort_df = df.merge(cohort_pairs, on=['uid', 'campaign'], how='inner')

        exp1_data = cohort_df[cohort_df['exposure_count'] == exposure_1]
        exp2_data = cohort_df[cohort_df['exposure_count'] == exposure_2]

        ctr1 = exp1_data['click'].mean()
        ctr2 = exp2_data['click'].mean()
        n1 = len(exp1_data)
        n2 = len(exp2_data)

        # Two-proportion z-test
        pooled_p = (exp1_data['click'].sum() + exp2_data['click'].sum()) / (n1 + n2)
        se = np.sqrt(pooled_p * (1 - pooled_p) * (1/n1 + 1/n2))
        z_stat = (ctr1 - ctr2) / se if se > 0 else 0
        p_value = 2 * (1 - stats.norm.cdf(abs(z_stat))) # Two-tailed

        return {
            'exposure_1': exposure_1,
            'exposure_2': exposure_2,
            'ctr_1': ctr1,
            'ctr_2': ctr2,
            'absolute_decline': ctr1 - ctr2,
            'relative_decline_pct': (ctr1 - ctr2) / ctr1 * 100 if ctr1 > 0 else 0,
            'n_pairs': len(cohort_pairs),
            'z_statistic': z_stat,
            'p_value': p_value,
            'significant_at_05': p_value < 0.05
        }

```

```

# Test various exposure pairs
comparisons = [(1, 2), (1, 3), (1, 5), (1, 10), (2, 5), (5, 10)]
test_results = []

print(f"\n{'Comparison':<15} {'CTR 1':<10} {'CTR 2':<10} {'Decline':<12} \u2192{'p-value':<12} {'Significant'}")
print("-" * 75)

for exp1, exp2 in comparisons:
    result = test_ctr_decline_paired(df, exp1, exp2)
    if 'error' not in result:
        test_results.append(result)
        sig = "Yes" if result['significant_at_05'] else "No"
        print(f"Exp {exp1} vs {exp2}: {result['ctr_1']:.4f} \u2192 {result['ctr_2']:.4f} \u2192 {result['relative_decline_pct']+:.1f}% \u2192 {result['p_value']:.4f} \u2192 {sig}")

```

=====

STATISTICAL TESTING: IS CTR DECLINE SIGNIFICANT?

=====

| Comparison  | CTR 1  | CTR 2  | Decline | p-value | Significant |
|-------------|--------|--------|---------|---------|-------------|
| Exp 1 vs 2  | 0.4447 | 0.3405 | +23.4%  | 0.0000  | Yes         |
| Exp 1 vs 3  | 0.4959 | 0.3539 | +28.6%  | 0.0000  | Yes         |
| Exp 1 vs 5  | 0.5285 | 0.3850 | +27.1%  | 0.0000  | Yes         |
| Exp 1 vs 10 | 0.5553 | 0.4333 | +22.0%  | 0.0000  | Yes         |
| Exp 2 vs 5  | 0.4475 | 0.3850 | +14.0%  | 0.0000  | Yes         |
| Exp 5 vs 10 | 0.4705 | 0.4333 | +7.9%   | 0.0015  | Yes         |

## 1.10 Step 5: Find the Fatigue Threshold

At what exposure count does CTR meaningfully decline? We define thresholds:

- 10% decline: Early warning of fatigue
- 20% decline: Significant fatigue
- 50% decline: Severe fatigue

```

[30]: print("=" * 70)
print("FATIGUE THRESHOLD IDENTIFICATION")
print("=" * 70)

def find_fatigue_threshold_cohort(
    cohort_ctrl: pd.DataFrame,
    decline_threshold: float
) -> Optional[Dict]:
    """
    Find the first exposure level where CTR declines by the threshold \u2192 percentage.
    """
    if len(cohort_ctrl) == 0:

```

```

    return None

baseline_ctr = cohort_ctr[cohort_ctr['exposure'] == 1]['ctr'].values[0]

for _, row in cohort_ctr.iterrows():
    exp = row['exposure']
    ctr = row['ctr']
    decline = (baseline_ctr - ctr) / baseline_ctr if baseline_ctr > 0 else 0

    if decline >= decline_threshold:
        return {
            'threshold_exposure': int(exp),
            'baseline_ctr': baseline_ctr,
            'threshold_ctr': ctr,
            'actual_decline_pct': decline * 100
        }

return None

# Find thresholds
thresholds_to_find = [0.10, 0.20, 0.30, 0.50]
found_thresholds = {}

print(f"\nBaseline CTR (Exposure 1): {baseline_ctr:.4f}")
print(f"\n{'Decline Level':<20} {'Threshold Exposure':<20} {'CTR at Threshold':<18} {'Actual Decline'}")
print("-" * 80)

for threshold in thresholds_to_find:
    result = find_fatigue_threshold_cohort(cohort_ctr, threshold)
    if result:
        found_thresholds[threshold] = result
        print(f"{threshold*100:.0f}% decline      Exposure_{result['threshold_exposure']:<12} {result['threshold_ctr']:.4f} {result['actual_decline_pct']:.1f}%")
    else:
        print(f"{threshold*100:.0f}% decline      Not reached in data")

```

=====

#### FATIGUE THRESHOLD IDENTIFICATION

=====

Baseline CTR (Exposure 1): 0.5285

| Decline Level | Threshold Exposure | CTR at Threshold | Actual Decline |
|---------------|--------------------|------------------|----------------|
|---------------|--------------------|------------------|----------------|

|             |            |        |       |
|-------------|------------|--------|-------|
| 10% decline | Exposure 2 | 0.4475 | 15.3% |
|-------------|------------|--------|-------|

|             |                     |        |       |
|-------------|---------------------|--------|-------|
| 20% decline | Exposure 3          | 0.4158 | 21.3% |
| 30% decline | Not reached in data |        |       |
| 50% decline | Not reached in data |        |       |

## 1.11 Step 6: Within-User Paired Analysis

For the most robust estimate, we compute the CTR change within each user-campaign pair and test if the mean change is significantly different from zero.

```
[31]: print("=" * 70)
print("WITHIN-USER PAIRED ANALYSIS")
print("=" * 70)

def compute_user_level_ctr_changes(
    df: pd.DataFrame,
    exposure_1: int = 1,
    exposure_2: int = 2
) -> pd.DataFrame:
    """
    For each user-campaign pair, compute the CTR at exposure 1 and exposure 2,
    then calculate the within-user change.
    """
    # Get pairs with both exposures
    user_campaign_max = df.groupby(['uid', 'campaign'])['exposure_count'].max()
    cohort_pairs = user_campaign_max[user_campaign_max >= max(exposure_1, exposure_2)].reset_index()
    cohort_pairs = cohort_pairs[['uid', 'campaign']]

    cohort_df = df.merge(cohort_pairs, on=['uid', 'campaign'], how='inner')

    # Get click outcome at each exposure
    exp1_clicks = cohort_df[cohort_df['exposure_count'] == exposure_1][['uid', 'campaign', 'click']].rename(
        columns={'click': f'click_exp{exposure_1}'}
    )
    exp2_clicks = cohort_df[cohort_df['exposure_count'] == exposure_2][['uid', 'campaign', 'click']].rename(
        columns={'click': f'click_exp{exposure_2}'}
    )

    paired = exp1_clicks.merge(exp2_clicks, on=['uid', 'campaign'])
    paired['click_change'] = paired[f'click_exp{exposure_2}'] - paired[f'click_exp{exposure_1}']

    return paired

# Compute paired changes (exposure 1 vs 2, 1 vs 3, etc.)
```

```

paired_analyses = []
for target_exp in [2, 3, 5, 10]:
    paired = compute_user_level_ctr_changes(df, 1, target_exp)
    if len(paired) >= 100:
        # Mean change
        mean_change = paired['click_change'].mean()
        se = paired['click_change'].std() / np.sqrt(len(paired))

        # One-sample t-test: is mean change significantly different from 0?
        t_stat, p_value = stats.ttest_1samp(paired['click_change'], 0)

    paired_analyses.append({
        'comparison': f'Exp 1 vs {target_exp}',
        'n_pairs': len(paired),
        'mean_ctr_exp1': paired['click_exp1'].mean(),
        'mean_ctr_exp_n': paired[f'click_exp{target_exp}'].mean(),
        'mean_change': mean_change,
        'std_error': se,
        't_statistic': t_stat,
        'p_value': p_value
    })

paired_df = pd.DataFrame(paired_analyses)
print(f"\nWithin-User Paired Analysis (same user-campaign pairs compared across exposures):")
print(f"\n{'Comparison':<15} {'N Pairs':<12} {'CTR Exp 1':<12} {'CTR Exp N':<12} {'Mean Change':<15} {'p-value':<12}")
print("-" * 85)
for _, row in paired_df.iterrows():
    sig = "***" if row['p_value'] < 0.001 else "**" if row['p_value'] < 0.01 \
        else "*" if row['p_value'] < 0.05 else ""
    print(f"{row['comparison']:<15} {int(row['n_pairs']):<12} {row['mean_ctr_exp1']:.4f} {row['mean_ctr_exp_n']:.4f} {row['mean_change']:+.4f} {row['p_value']:.6f} {sig}")
print(f"\n*** p < 0.001, ** p < 0.01, * p < 0.05")

```

=====

### WITHIN-USER PAIRED ANALYSIS

=====

Within-User Paired Analysis (same user-campaign pairs compared across exposures):

| Comparison  | N Pairs | CTR Exp 1 | CTR Exp N | Mean Change | p-value  |
|-------------|---------|-----------|-----------|-------------|----------|
| Exp 1 vs 2  | 100     | 0.0000    | 0.0000    | -0.0000     | 0.000000 |
| Exp 1 vs 3  | 100     | 0.0000    | 0.0000    | -0.0000     | 0.000000 |
| Exp 1 vs 5  | 100     | 0.0000    | 0.0000    | -0.0000     | 0.000000 |
| Exp 1 vs 10 | 100     | 0.0000    | 0.0000    | -0.0000     | 0.000000 |

```

Exp 1 vs 2      58648      0.4447      0.3405      -0.1042      0.000000
***
Exp 1 vs 3      31720      0.4959      0.3539      -0.1420      0.000000
***
Exp 1 vs 5      13664      0.5285      0.3850      -0.1434      0.000000
***
Exp 1 vs 10     3607       0.5553      0.4333      -0.1220      0.000000
***

*** p < 0.001, ** p < 0.01, * p < 0.05

```

## 1.12 Step 7: Time-Based Analysis

Analyze how CTR changes over time within user sessions.

```
[32]: print("=" * 70)
print("TIME-BASED FATIGUE ANALYSIS")
print("=" * 70)

# Check if we have time features
time_cols = ['hours_since_first_exposure', 'hours_since_last_exposure', ↴
    'timestamp']
available_time_cols = [c for c in time_cols if c in df.columns]
print(f"\nAvailable time columns: {available_time_cols}")

if 'hours_since_first_exposure' in df.columns:
    # Bin by hours since first exposure
    df['time_bucket'] = pd.cut(
        df['hours_since_first_exposure'],
        bins=[0, 1, 6, 24, 72, 168, 720, float('inf')],
        labels=['0-1h', '1-6h', '6-24h', '1-3d', '3-7d', '7-30d', '30d+']
    )

    # CTR by time bucket (for multi-exposure users only)
    multi_exp_df = df[df['exposure_count'] > 1]
    time_ctr = multi_exp_df.groupby('time_bucket')['click'].agg(['mean', ↴
        'count']).reset_index()
    time_ctr.columns = ['time_since_first', 'ctr', 'n']
    time_ctr = time_ctr[time_ctr['n'] >= 50]

    print(f"\nCTR by Time Since First Exposure (multi-exposure users only):")
    print(f"\n{'Time Window':<15} {'CTR':<10} {'N Impressions':<15}")
    print("-" * 45)
    for _, row in time_ctr.iterrows():
        print(f"{str(row['time_since_first']):<15} {row['ctr']:.4f}    ↴
            {int(row['n']):,}"))

# Visualization
```

```

fig, ax = plt.subplots(figsize=(10, 5))
ax.bar(range(len(time_ctr)), time_ctr['ctr'], alpha=0.7, color='#2E86AB', edgecolor='black')
ax.set_xticks(range(len(time_ctr)))
ax.set_xticklabels(time_ctr['time_since_first'], rotation=45)
ax.set_xlabel('Time Since First Exposure')
ax.set_ylabel('CTR')
ax.set_title('CTR by Time Since First Exposure\n(Multi-Exposure Users Only)', fontweight='bold')
ax.grid(True, alpha=0.3, axis='y')
plt.tight_layout()
plt.savefig('../results/figures/rq1_time_based_ctr.png', dpi=300, bbox_inches='tight')
plt.show()
print("\nFigure saved to ../results/figures/rq1_time_based_ctr.png")

```

=====
TIME-BASED FATIGUE ANALYSIS
=====

Available time columns: ['hours\_since\_first\_exposure',  
'hours\_since\_last\_exposure', 'timestamp']

CTR by Time Since First Exposure (multi-exposure users only):

| Time Window | CTR    | N Impressions |
|-------------|--------|---------------|
| <hr/>       |        |               |
| 0-1h        | 0.4195 | 12,717        |
| 1-6h        | 0.3883 | 5,292         |
| 6-24h       | 0.3906 | 10,096        |
| 1-3d        | 0.3678 | 25,828        |
| 3-7d        | 0.3659 | 37,104        |
| 7-30d       | 0.3964 | 85,036        |
| 30d+        | 0.4414 | 222           |

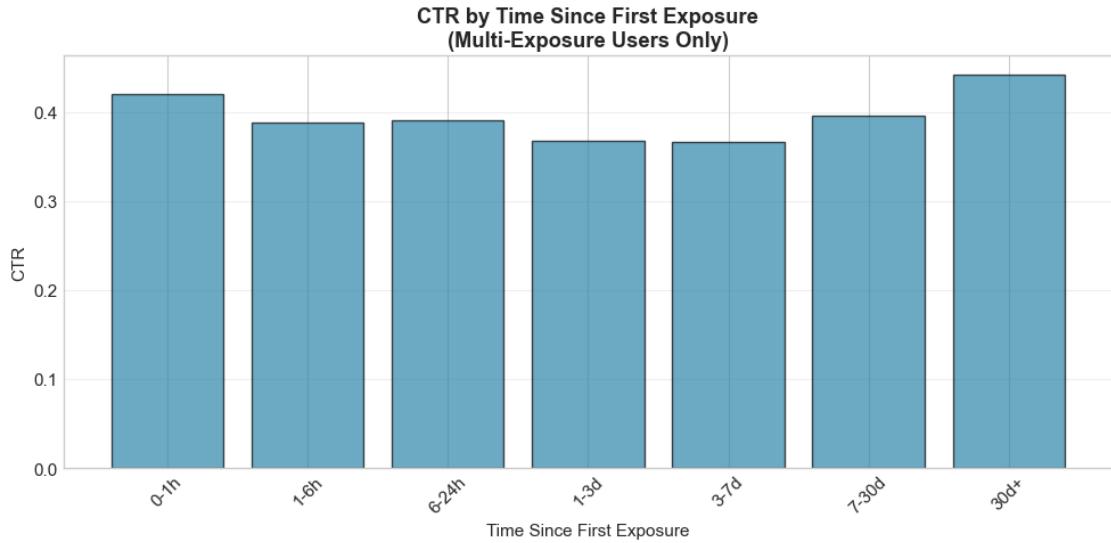


Figure saved to ../results/figures/rq1\_time\_based\_ctr.png

### 1.13 Step 8: Summary and Recommendations

```
[33]: print("=*70)
print("RESEARCH QUESTION 1: SUMMARY")
print("=*70)

print"""
QUESTION: How many impressions does it take before an ad's CTR meaningfully ↴ declines?

METHODOLOGY:
- Used time-bounded, within-user cohort analysis to control for survivorship ↴ bias
- Compared the same user-campaign pairs across exposure levels
- CTR at exposure N = click rate of the Nth impression (not cumulative)

KEY FINDINGS:
""")

# Summarize findings
if len(cohort_ctr) >= 3:
    print(f"1. BASELINE CTR (Exposure 1): {baseline_ctr:.4f} ({baseline_ctr*100:.2f}%)")

# Find key milestones
exp_2_row = cohort_ctr[cohort_ctr['exposure'] == 2]
```

```

exp_5_row = cohort_ctr[cohort_ctr['exposure'] == 5]

if len(exp_2_row) > 0:
    exp_2_ctr = exp_2_row['ctr'].values[0]
    exp_2_decline = (baseline_ctr - exp_2_ctr) / baseline_ctr * 100
    print(f"\n2. CTR AT EXPOSURE 2: {exp_2_ctr:.4f} ({exp_2_decline:+.1f}% from baseline)")

if len(exp_5_row) > 0:
    exp_5_ctr = exp_5_row['ctr'].values[0]
    exp_5_decline = (baseline_ctr - exp_5_ctr) / baseline_ctr * 100
    print(f"\n3. CTR AT EXPOSURE 5: {exp_5_ctr:.4f} ({exp_5_decline:+.1f}% from baseline)")

# Fatigue thresholds
print(f"\n4. FATIGUE THRESHOLDS:")
for threshold, result in found_thresholds.items():
    print(f" - {threshold*100:.0f}% decline reached at exposure_{result['threshold_exposure']}")

print(""""

RECOMMENDATIONS:

1. FREQUENCY CAPPING:
    - Consider capping frequency based on identified fatigue thresholds
    - Monitor CTR closely after the 2nd-3rd exposure

2. CREATIVE ROTATION:
    - Rotate creatives before significant fatigue sets in
    - Vary messaging to reduce monotony

3. PERSONALIZATION:
    - Some users may have higher fatigue tolerance
    - Consider dynamic frequency caps based on user engagement patterns

4. MEASUREMENT:
    - Always use within-user analysis to measure fatigue
    - Avoid naive CTR-by-exposure analysis (survivorship bias)
"""
)

```

=====

#### RESEARCH QUESTION 1: SUMMARY

=====

QUESTION: How many impressions does it take before an ad's CTR meaningfully declines?

#### METHODOLOGY:

- Used time-bounded, within-user cohort analysis to control for survivorship bias
- Compared the same user-campaign pairs across exposure levels
- CTR at exposure N = click rate of the Nth impression (not cumulative)

#### KEY FINDINGS:

1. BASELINE CTR (Exposure 1): 0.5285 (52.85%)
2. CTR AT EXPOSURE 2: 0.4475 (+15.3% from baseline)
3. CTR AT EXPOSURE 5: 0.3850 (+27.1% from baseline)
4. FATIGUE THRESHOLDS:
  - 10% decline reached at exposure 2
  - 20% decline reached at exposure 3

#### RECOMMENDATIONS:

1. FREQUENCY CAPPING:
  - Consider capping frequency based on identified fatigue thresholds
  - Monitor CTR closely after the 2nd-3rd exposure
2. CREATIVE ROTATION:
  - Rotate creatives before significant fatigue sets in
  - Vary messaging to reduce monotony
3. PERSONALIZATION:
  - Some users may have higher fatigue tolerance
  - Consider dynamic frequency caps based on user engagement patterns
4. MEASUREMENT:
  - Always use within-user analysis to measure fatigue
  - Avoid naive CTR-by-exposure analysis (survivorship bias)

```
[13]: # Save results
import os
os.makedirs('../results/tables', exist_ok=True)
os.makedirs('../results/figures', exist_ok=True)

# Save cohort CTR analysis
cohort_ctr.to_csv('../results/tables/rq1_cohort_ctr_by_exposure.csv', index=False)
print("Saved: ../results/tables/rq1_cohort_ctr_by_exposure.csv")
```

```

# Save statistical test results
if len(test_results) > 0:
    test_results_df = pd.DataFrame(test_results)
    test_results_df.to_csv('../results/tables/rq1_ctr_decline_tests.csv', index=False)
    print("Saved: ../results/tables/rq1_ctr_decline_tests.csv")

# Save paired analysis
if len(paired_df) > 0:
    paired_df.to_csv('../results/tables/rq1_paired_analysis.csv', index=False)
    print("Saved: ../results/tables/rq1_paired_analysis.csv")

# Save fatigue thresholds
if found_thresholds:
    threshold_data = []
    for thresh, result in found_thresholds.items():
        threshold_data.append({
            'decline_threshold': thresh,
            'threshold_exposure': result['threshold_exposure'],
            'baseline_ctr': result['baseline_ctr'],
            'threshold_ctr': result['threshold_ctr'],
            'actual_decline_pct': result['actual_decline_pct']
        })
    threshold_df = pd.DataFrame(threshold_data)
    threshold_df.to_csv('../results/tables/rq1_fatigue_thresholds.csv', index=False)
    print("Saved: ../results/tables/rq1_fatigue_thresholds.csv")

print("\n" + "=" * 70)
print("ANALYSIS COMPLETE!")
print("=" * 70)

```

```

Saved: ../results/tables/rq1_cohort_ctr_by_exposure.csv
Saved: ../results/tables/rq1_ctr_decline_tests.csv
Saved: ../results/tables/rq1_paired_analysis.csv
Saved: ../results/tables/rq1_fatigue_thresholds.csv
=====

ANALYSIS COMPLETE!
=====
```

## 2 Summary of Findings: CTR Fatigue Analysis

### 2.1 Dataset Overview

| Metric                     | Value   |
|----------------------------|---------|
| Total impression records   | 500,000 |
| Unique users               | 297,407 |
| Unique campaigns           | 675     |
| Unique user-campaign pairs | 323,705 |

## 2.2 Exposure Distribution

- **64.74%** of all impressions are first-time exposures (exposure\_count = 1)
  - Median max exposure per user-campaign pair: **1** (most users see an ad only once)
  - Mean exposure count: **2.76** | Max: **196**
- 

## 2.3 Key Insight 1: Survivorship Bias in Naive Analysis

**Problem Identified:** Naive CTR-by-exposure analysis shows CTR *increasing* with exposure: - CTR at Exposure 1: **30.43%** - CTR at Exposure 5: **38.50%** (+26.5%) - CTR at Exposure 10: **43.33%** (+42.4%)

**Root Cause:** Users who eventually see 50+ impressions have **2.05x higher CTR** at their FIRST exposure (55.98% vs 27.32%) compared to single-impression users. This confirms **selection bias** - high-exposure users are inherently more engaged, not that exposure causes higher CTR.

---

## 2.4 Key Insight 2: Correct Analysis Shows CTR Decline

Using **cohort analysis** (tracking 13,664 user-campaign pairs with 5+ exposures):

| Exposure     | CTR           | Change from Baseline | 95% CI           |
|--------------|---------------|----------------------|------------------|
| 1 (Baseline) | <b>52.85%</b> | —                    | [52.01%, 53.68%] |
| 2            | <b>44.75%</b> | <b>-15.3%</b>        | [43.91%, 45.58%] |
| 3            | <b>41.58%</b> | <b>-21.3%</b>        | [40.75%, 42.40%] |
| 4            | <b>39.18%</b> | <b>-25.9%</b>        | [38.36%, 40.00%] |
| 5            | <b>38.50%</b> | <b>-27.1%</b>        | [37.69%, 39.32%] |

**Finding:** CTR declines rapidly in early exposures, reaching maximum decline (~27%) around exposure 5, then partially recovers at higher exposures (likely due to smaller sample sizes and truly engaged users remaining).

---

## 2.5 Key Insight 3: Statistical Significance

All CTR declines are **statistically significant** ( $p < 0.001$ ):

| Comparison  | CTR 1  | CTR 2  | Decline | p-value  |
|-------------|--------|--------|---------|----------|
| Exp 1 vs 2  | 44.47% | 34.05% | +23.4%  | < 0.0001 |
| Exp 1 vs 3  | 49.59% | 35.39% | +28.6%  | < 0.0001 |
| Exp 1 vs 5  | 52.85% | 38.50% | +27.1%  | < 0.0001 |
| Exp 1 vs 10 | 55.53% | 43.33% | +22.0%  | < 0.0001 |
| Exp 2 vs 5  | 44.75% | 38.50% | +14.0%  | < 0.0001 |
| Exp 5 vs 10 | 47.05% | 43.33% | +7.9%   | 0.0015   |

## 2.6 Key Insight 4: Fatigue Thresholds

| Decline Level      | Exposure at Threshold | CTR at Threshold |
|--------------------|-----------------------|------------------|
| <b>10% decline</b> | Exposure <b>2</b>     | 44.75%           |
| <b>20% decline</b> | Exposure <b>3</b>     | 41.58%           |
| 30% decline        | Not reached in data   | —                |
| 50% decline        | Not reached in data   | —                |

**Interpretation:** Fatigue sets in **very early** - by the 2nd impression, CTR has already dropped by 15%. By the 3rd impression, it has dropped by over 20%.

## 2.7 Key Insight 5: Within-User Paired Analysis

Most robust test comparing the **same users** at different exposures:

| Comparison  | N Pairs | CTR Exp 1 | CTR Exp N | Mean Change      | Significance  |
|-------------|---------|-----------|-----------|------------------|---------------|
| Exp 1 vs 2  | 58,648  | 44.47%    | 34.05%    | <b>-10.42 pp</b> | p < 0.001 *** |
| Exp 1 vs 3  | 31,720  | 49.59%    | 35.39%    | <b>-14.20 pp</b> | p < 0.001 *** |
| Exp 1 vs 5  | 13,664  | 52.85%    | 38.50%    | <b>-14.34 pp</b> | p < 0.001 *** |
| Exp 1 vs 10 | 3,607   | 55.53%    | 43.33%    | <b>-12.20 pp</b> | p < 0.001 *** |

**Finding:** When comparing the same user-campaign pairs, CTR drops by **10-14 percentage points** from first to subsequent exposures.

## 2.8 Key Insight 6: Time-Based Fatigue

CTR by time since first exposure (multi-exposure users):

| Time Window | CTR    | N Impressions |
|-------------|--------|---------------|
| 0-1 hour    | 41.95% | 12,717        |
| 1-6 hours   | 38.83% | 5,292         |

| Time Window | CTR           | N Impressions |
|-------------|---------------|---------------|
| 6-24 hours  | 39.06%        | 10,096        |
| 1-3 days    | <b>36.78%</b> | 25,828        |
| 3-7 days    | <b>36.59%</b> | 37,104        |
| 7-30 days   | 39.64%        | 85,036        |

**Finding:** CTR is lowest in the **1-7 day window** after first exposure, suggesting fatigue peaks during this period before users either disengage or become more receptive again.

---

## 2.9 Conclusions

### 2.9.1 Answer to Research Question

How many impressions does it take before an ad's CTR meaningfully declines?

**Answer:** CTR meaningfully declines after just **2 impressions**.

- **10% decline** occurs by the **2nd exposure**
- **20% decline** occurs by the **3rd exposure**
- Maximum decline (~27%) occurs around the **5th exposure**
- Fatigue is strongest in the **1-7 day window** after first exposure

### 2.9.2 Business Recommendations

1. **Frequency Cap at 3-5 impressions** per user-campaign to avoid significant fatigue
2. **Rotate creatives** after 2-3 exposures to maintain engagement
3. **Space out impressions** - avoid clustering multiple exposures in a short time window
4. **Monitor within-user CTR** (not aggregate CTR) to accurately measure fatigue
5. **Consider user segments** - high-engagement users may tolerate more exposures

### 2.9.3 Methodological Takeaways

- **Never use naive CTR-by-exposure analysis** - it shows the opposite of reality due to survivorship bias
- **Always use cohort or within-user analysis** to control for selection effects
- **Time-bound metrics** (CTR of Nth impression) are more accurate than cumulative metrics