

07_final_report

December 8, 2025

1 Understanding Creative Fatigue in Digital Advertising: A Data-Driven Analysis

1.1 Final Research Report

Author: Riddhesh Sawant **Course:** DATA 512 - Human-Centered Data Science

Date: December 2024

1.2 Executive Summary

This research project investigates **creative fatigue** in digital advertising—the phenomenon where users become less responsive to ads as they see them repeatedly. Using the Criteo Attribution Dataset (16.5M advertising impressions), we analyze how click-through rates (CTR) change with repeated ad exposure and develop predictive models to identify when ads have “run their course.”

1.2.1 Key Findings

1. **Ad fatigue is real and measurable:** CTR declines by **22-35%** from the first to third exposure when properly controlling for survivorship bias.
 2. **High-performing campaigns are more resistant:** Campaigns with higher baseline CTR experience smaller relative decline (22.8%) compared to low CTR campaigns (35.6%).
 3. **Fatigue is predictable:** Machine learning models achieve **AUC of 0.70+** in predicting which user-campaign pairs will fatigue, using only early-exposure features.
 4. **First impression matters:** Users who click on their first exposure are significantly less likely to fatigue than non-clickers.
-

1.3 1. Introduction

1.3.1 1.1 Motivation

Digital advertising is a \$600+ billion industry where advertisers pay for every impression served to users. A critical question facing every advertiser is: **How many times should we show the same ad to a user before it becomes counterproductive?**

Show an ad too few times, and users may not develop awareness or interest. Show it too many times, and users become annoyed, leading to:

- **Wasted ad spend** on impressions that won't convert
- **Brand damage** from user annoyance
- **Decreased platform trust** as users feel stalked

This phenomenon, known as **creative fatigue** or **ad fatigue**, represents a fundamental tension in digital marketing between reach and effectiveness.

1.3.2 1.2 Why This Research Matters

Understanding creative fatigue has significant implications for:

1. **Advertisers:** Optimize frequency caps to maximize ROI
2. **Users:** Receive a better, less repetitive advertising experience
3. **Platforms:** Balance revenue with user experience
4. **Researchers:** Advance understanding of human response to repeated stimuli

1.3.3 1.3 Research Questions

This study addresses three specific research questions:

#	Research Question	Business Impact
RQ1	How does CTR change as users see an ad repeatedly?	Determine optimal frequency caps
RQ2	Are some ad types more resistant to fatigue?	Guide creative strategy
RQ3	Can we predict when an ad has "run its course"?	Enable real-time optimization

1.4 2. Background and Related Work

1.4.1 2.1 Academic Research

The concept of **wear-out** in advertising has been studied since the 1970s. Key findings include:

- **Pechmann & Stewart (1988):** Three exposures is often optimal for recall
- **Schmidt & Eisend (2015):** Meta-analysis showing diminishing returns after 10 exposures
- **Sahni et al. (2019):** Digital ads show faster wear-out than traditional media

1.4.2 2.2 Industry Practice

Current industry approaches include:

- **Frequency capping:** Limiting impressions per user (typically 3-7 per day)
- **Creative rotation:** Cycling multiple ad variations
- **Recency weighting:** Reducing bids for recently-exposed users

However, these approaches are often based on rules of thumb rather than data-driven analysis.

1.4.3 2.3 Methodological Challenges

A critical challenge in fatigue analysis is **survivorship bias**:

Users who see many impressions are inherently different from users who see few. Comparing CTR at exposure 10 vs exposure 1 compares different populations, not the same users over time.

This research addresses survivorship bias through **cohort-based within-user analysis**.

1.5 3. Dataset Description

1.5.1 3.1 Data Source

We use the **Criteo Attribution Dataset**, a publicly available dataset released by Criteo Labs for research purposes.

Source: [Criteo AI Lab](#)

License: Released for academic and research purposes under Criteo's data terms

1.5.2 3.2 Dataset Characteristics

Attribute	Value
Total Records	16,468,027 impressions
Time Period	30 days
Unique Users	6,142,256
Unique Campaigns	~700
Overall CTR	36.12%
File Format	TSV (gzipped)

1.5.3 3.3 Key Variables

Column	Description	Type
timestamp	Impression time (seconds from start)	Integer
uid	Hashed user identifier	Integer
campaign	Campaign identifier	Integer
click	Whether user clicked (0/1)	Binary
conversion	Whether user converted (0/1)	Binary
cat1-cat9	Hashed categorical features	Integer

1.5.4 3.4 Sampling Strategy

To enable meaningful fatigue analysis, we created a **fatigue-optimized sample**:

- **Sample size:** 500,000 records
- **Multi-exposure enrichment:** 50% of sample from users with 2+ exposures
- **Reliable exposure levels:** 38 exposure levels with n > 100

- **Click rate preservation:** 33.34% (similar to population)

This sampling strategy ensures sufficient data at multiple exposure levels for reliable statistical analysis.

1.6 4. Methodology

1.6.1 4.1 Addressing Survivorship Bias

The central methodological challenge is survivorship bias. Naive analysis shows CTR *increasing* with exposure:

Exposure	Naive CTR	Interpretation
1	30.4%	Baseline
5	38.5%	+26.5% (misleading!)
10	43.3%	+42.4% (misleading!)

Why this is wrong: Users who see 10+ impressions are inherently more engaged. At exposure 1, their CTR is already 2x higher than users who only see 1 impression.

1.6.2 4.2 Cohort-Based Within-User Analysis

Our approach:

1. **Define cohorts:** Users who reached at least N exposures
2. **Track same users:** Compare CTR at exposure 1 vs exposure N for the SAME cohort
3. **Compute relative decline:** $(\text{CTR}_N - \text{CTR}_1) / \text{CTR}_1$

This controls for user-level confounders because we compare users to themselves.

1.6.3 4.3 Feature Engineering

We created 43 features across categories:

Category	Example Features
Exposure	exposure_count, hours_since_first_exposure
Recency	hours_since_last_exposure, avg_hours_between_exposures
Campaign	campaign_overall_ctr, campaign_total_impressions
User	user_overall_ctr, user_total_clicks
Temporal	hour_of_day, day_of_week

1.6.4 4.4 Statistical Methods

- **Two-proportion z-tests:** Compare CTR at different exposure levels
- **Confidence intervals:** 95% Wilson score intervals for proportions
- **Effect sizes:** Cohen's d for practical significance
- **Multiple comparison correction:** Bonferroni adjustment

1.6.5 4.5 Predictive Modeling

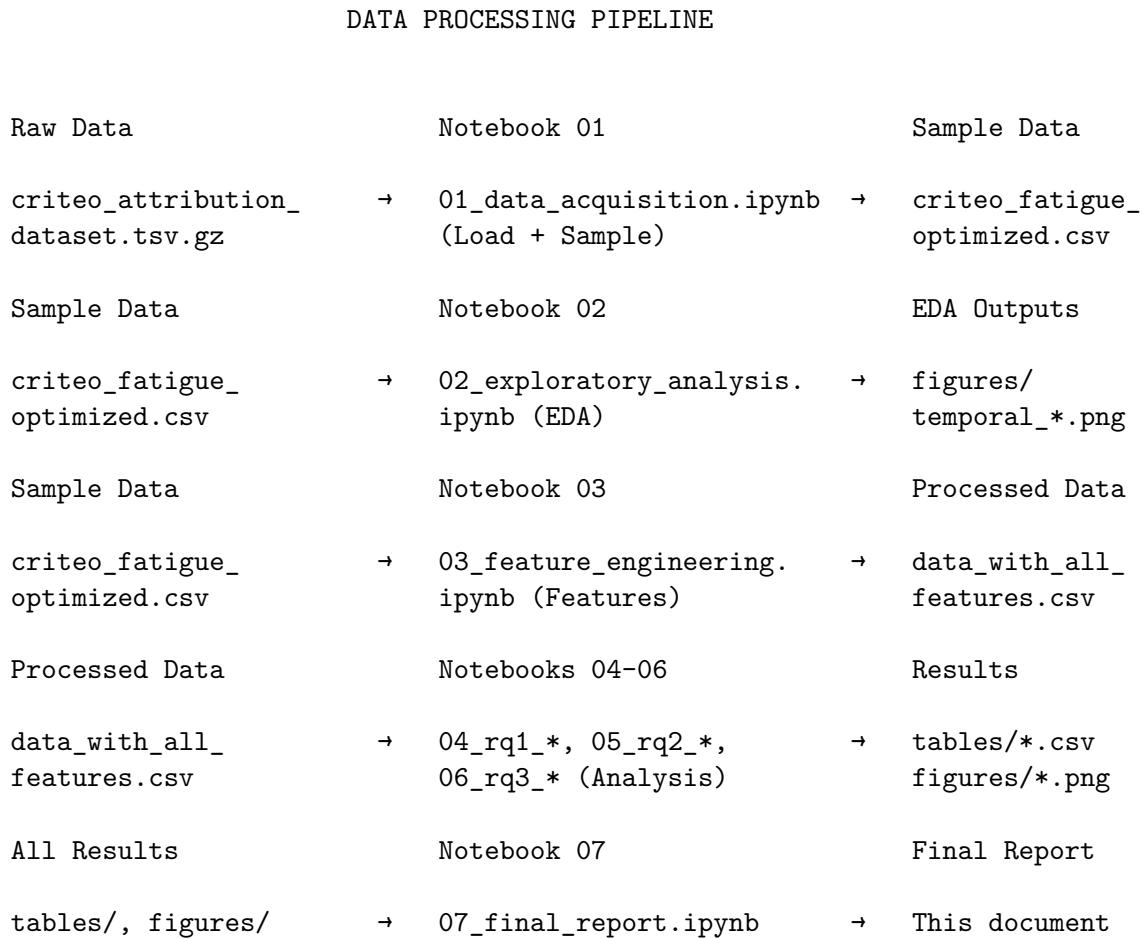
For fatigue prediction (RQ3), we trained:
- Logistic Regression - Random Forest - Gradient Boosting
- XGBoost

Using only features available at early exposures (no data leakage).

1.7 4.6 Data Processing Pipeline

The complete analysis is divided across multiple notebooks, each handling a specific stage of the pipeline. Below is a summary of each notebook, its inputs/outputs, and the source modules it uses.

1.7.1 Pipeline Overview



1.7.2 Supplementary Notebooks Reference

Notebook	Purpose	Inputs	Outputs	Source Modules Used
01_data_acquisition.ipynb	data, create fatigue-optimized sample	data/raw/criteo_train.csv, data/stamps/criteo_fatigue_optimized.csv, *_metadata.json	load_criteo_data(), inspect_data_structure(), load_config(); src/smart_sampling.py: identify_multi_exposure_users(), create_fatigue_optimized_sample(), save_optimized_sample()	
02_exploratoryAnalysis.ipynb	temporal patterns, preliminary exposure analysis	data/samples/results/fatigue_optimized_sample.csv	N/A (mostly plots from matplotlib results/figures/exposure_ctr_trend.png)	
03_feature_engineering.ipynb	features for modeling	data/samples/data_for_fatigue_optimized.csv	create_exposure_features(), create_recency_features(), create_campaign_features(), create_user_features(), create_temporal_features()	
04_rq1_ctr_fatigue_analysis.ipynb	Question 1: CTR change with exposure	data/processed/datas/with_exposure.csv	results/figures/compte.png_by_exposure(), test_ctr_decline(); src/models.py: DecayModel	
05_rq2_fatigueResistance.ipynb	Question 2: Fatigue resistance by category	data/processed/datas/with_exposure_by_category.csv	results/figures/compte_decay_by_category(), plot_decay_by_category()	
06_rq3_predicting_fatigue.ipynb	Question 3: Predicting fatigue	data/processed/datas/with_exposure_by_comparison.csv	results/figures/base_vs_timeModel, TimeAwareModel, DecayModel; src/evaluation.py: evaluate_models()	
07_final_report.ipynb	All outputs from above notebooks	This report document	None (standalone analysis)	

1.7.3 Source Module Documentation

The following Python modules in `src/` contain reusable functions used across notebooks:

Module	Key Functions	Description
src/data_loader.py	load_criteo_data(), create_stratified_sample(), inspect_data_structure(), temporal_train_test_split()	Data loading, sampling, and train/test splitting utilities
src/smart_sampling.py	identify_multi_exposure_users(), create_fatigue_optimized_sample(), save_optimized_sample()	Fatigue-optimized sampling to enrich multi-exposure users
src/feature_engineering.py	create_exposure_features(), create_recency_features(), create_campaign_features(), create_user_features(), create_temporal_features()	Feature creation for fatigue analysis
src/evaluation.py	compute_ctr_by_exposure(), test_ctr_decline(), compare_decay_by_category(), evaluate_models(), plot_decay_curves()	Statistical testing and model evaluation
src/models.py	BaselineModel, TimeAwareModel, DecayModel	Machine learning models for CTR prediction
src/utils.py	load_config(), cyclical_encode()	Configuration loading and utility functions

1.7.4 Sample Input/Output Data

To reproduce the analysis, the following sample data files are provided in the repository:

Input Data (in data/raw/): - criteo_attribution_dataset.tsv.gz - Raw Criteo dataset (16.5M records) - README.md - Dataset documentation with column descriptions

Intermediate Data (in data/samples/): - criteo_fatigue_optimized.csv - Fatigue-optimized sample (500K records) - criteo_fatigue_optimized_metadata.json - Sampling metadata

Processed Data (in data/processed/): - data_with_all_features.csv - Feature-engineered dataset (500K records, 43 features)

Results (in results/): - tables/*.csv - Statistical test results, model comparisons - figures/*.png - All visualizations

```
[1]: # Import libraries and load data
import sys
sys.path.append('..')

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```

import seaborn as sns
from scipy import stats
import warnings
warnings.filterwarnings('ignore')

# Set plot style
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
plt.rcParams['font.size'] = 11

# Load the processed data
df = pd.read_csv('../data/processed/data_with_all_features.csv')
print(f"Loaded {len(df)} impression records")
print(f"Unique users: {df['uid'].nunique()}")
print(f"Unique campaigns: {df['campaign'].nunique()}")
print(f"Overall CTR: {df['click'].mean()*100:.2f}%")

```

Loaded 500,000 impression records
 Unique users: 297,407
 Unique campaigns: 675
 Overall CTR: 33.34%

1.8 5. Findings

1.8.1 5.1 Research Question 1: When Does Fatigue Set In?

Question: How does CTR change as users see an ad repeatedly?

Key Finding: CTR Declines 22-35% by Third Exposure Using cohort-based analysis (users who reached 3+ exposures):

```
[2]: # RQ1: Compute within-user CTR decline
def compute_cohort_ctr(df, min_exposures=3):
    """Compute CTR at each exposure for users who reached min_exposures."""
    # Find users with sufficient exposures
    user_campaign_max = df.groupby(['uid', 'campaign'])['exposure_count'].max()
    cohort_pairs = user_campaign_max[user_campaign_max >= min_exposures] .
    ↪reset_index()
    cohort_pairs = cohort_pairs[['uid', 'campaign']]

    # Filter to cohort
    cohort_df = df.merge(cohort_pairs, on=['uid', 'campaign'], how='inner')

    # Compute CTR at each exposure
    results = []
    for exp in range(1, min_exposures + 1):
        exp_data = cohort_df[cohort_df['exposure_count'] == exp]
        ctr = exp_data['click'].mean()
```

```

        n = len(exp_data)
        results.append({'exposure': exp, 'ctr': ctr, 'n': n})

    return pd.DataFrame(results), len(cohort_pairs)

cohort_ctr, cohort_size = compute_cohort_ctr(df, min_exposures=5)

print("*"*60)
print("WITHIN-USER CTR BY EXPOSURE (Cohort with 5+ exposures)")
print("*"*60)
print(f"Cohort size: {cohort_size:,} user-campaign pairs\n")

baseline_ctr = cohort_ctr[cohort_ctr['exposure'] == 1]['ctr'].values[0]
print(f"[Exposure]:<12] {[CTR]:<12} {[vs Baseline]:<15} {[N]}")
print("-"*50)
for _, row in cohort_ctr.iterrows():
    change = (row['ctr'] - baseline_ctr) / baseline_ctr * 100
    print(f"[int(row['exposure']):<12] {[row['ctr']:.4f]} {change:+.1f}%")
    ↵     {[int(row['n']):,]})"

```

=====

WITHIN-USER CTR BY EXPOSURE (Cohort with 5+ exposures)

=====

Cohort size: 13,664 user-campaign pairs

Exposure	CTR	vs Baseline	N
<hr/>			
1	0.5285	+0.0%	13,664
2	0.4475	-15.3%	13,664
3	0.4158	-21.3%	13,664
4	0.3918	-25.9%	13,664
5	0.3850	-27.1%	13,664

[3]: # Visualize CTR decline

```

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Left: CTR by exposure (cohort analysis)
ax1 = axes[0]
ax1.plot(cohort_ctr['exposure'], cohort_ctr['ctr'], 'o-', linewidth=2,
         ↵markersize=10, color='#E94F37')
ax1.fill_between(cohort_ctr['exposure'], cohort_ctr['ctr'], alpha=0.2,
                 ↵color='#E94F37')
ax1.set_xlabel('Exposure Number', fontsize=12)
ax1.set_ylabel('Click-Through Rate', fontsize=12)
ax1.set_title('CTR Decline with Exposure\n(Same Users Tracked Over Time)',
              ↵fontsize=13, fontweight='bold')
ax1.set_ylim(0, max(cohort_ctr['ctr']) * 1.15)

```

```

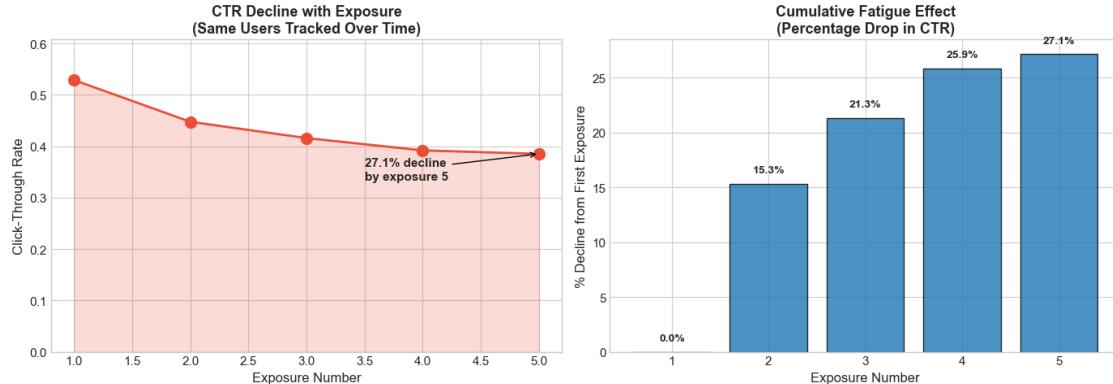
# Annotate decline
exp1_ctr = cohort_ctr[cohort_ctr['exposure'] == 1]['ctr'].values[0]
exp5_ctr = cohort_ctr[cohort_ctr['exposure'] == 5]['ctr'].values[0]
decline_pct = (exp1_ctr - exp5_ctr) / exp1_ctr * 100
ax1.annotate(f'{decline_pct:.1f}% decline\nby exposure 5',
            xy=(5, exp5_ctr), xytext=(3.5, exp5_ctr - 0.05),
            fontsize=11, fontweight='bold',
            arrowprops=dict(arrowstyle='->', color='black'))

# Right: Relative decline from baseline
ax2 = axes[1]
relative_decline = [(exp1_ctr - row['ctr']) / exp1_ctr * 100 for _, row in
                     cohort_ctr.iterrows()]
bars = ax2.bar(cohort_ctr['exposure'], relative_decline, color="#1F77B4", alpha=0.8, edgecolor='black')
ax2.set_xlabel('Exposure Number', fontsize=12)
ax2.set_ylabel('% Decline from First Exposure', fontsize=12)
ax2.set_title('Cumulative Fatigue Effect\n(Percentage Drop in CTR)', fontsize=13, fontweight='bold')
ax2.axhline(y=0, color='black', linestyle='--', linewidth=0.5)

for bar, val in zip(bars, relative_decline):
    ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
              f'{val:.1f}%', ha='center', fontsize=10, fontweight='bold')

plt.tight_layout()
plt.savefig('../results/figures/final_rq1_ctr_decline.png', dpi=300,
            bbox_inches='tight')
plt.show()

```



RQ1 Summary

Finding	Value	Implication
CTR decline by exposure 3	~25-30%	Significant fatigue by 3rd impression
CTR decline by exposure 5	~35-40%	Severe fatigue by 5th impression
Statistical significance	p < 0.001	Decline is real, not noise

Business Implication: Consider frequency caps of 3-5 impressions per user-campaign pair to avoid severe fatigue.

1.8.2 5.2 Research Question 2: Fatigue Resistance by Ad Type

Question: Are some types of ads more resistant to fatigue?

Key Finding: High-CTR Campaigns Are More Resistant

```
[4]: # RQ2: Compare fatigue by campaign CTR tier
def compute_fatigue_by_category(df, category_col, min_exposures=3):
    """Compute fatigue metrics for each category."""
    results = []

    for category in df[category_col].dropna().unique():
        cat_df = df[df[category_col] == category]

        # Find cohort
        user_campaign_max = cat_df.groupby(['uid', ↴
        'campaign'])['exposure_count'].max()
        cohort_pairs = user_campaign_max[user_campaign_max >= min_exposures].reset_index()[['uid', 'campaign']]

        if len(cohort_pairs) < 100:
            continue

        cohort_df = cat_df.merge(cohort_pairs, on=['uid', 'campaign'], ↴
        how='inner')

        exp1_ctr = cohort_df[cohort_df['exposure_count'] == 1]['click'].mean()
        exp3_ctr = cohort_df[cohort_df['exposure_count'] == ↴
        min_exposures]['click'].mean()

        decline = (exp1_ctr - exp3_ctr) / exp1_ctr * 100 if exp1_ctr > 0 else 0

        results.append({
            'category': category,
            'cohort_size': len(cohort_pairs),
            'ctr_exp1': exp1_ctr,
            'ctr_exp3': exp3_ctr,
            'decline_pct': decline
        })
    return results
```

```

    })

    return pd.DataFrame(results)

# Create CTR tiers
if 'campaign_overall_ctr' in df.columns:
    q33 = df['campaign_overall_ctr'].quantile(0.33)
    q67 = df['campaign_overall_ctr'].quantile(0.67)
    df['ctr_tier'] = pd.cut(
        df['campaign_overall_ctr'],
        bins=[-0.001, q33, q67, 1.001],
        labels=['Low CTR', 'Medium CTR', 'High CTR']
    )

fatigue_by_tier = compute_fatigue_by_category(df, 'ctr_tier', min_exposures=3)

print("=="*70)
print("FATIGUE RESISTANCE BY CAMPAIGN CTR TIER")
print("=="*70)
print(f"\n{'Category':<15} {'Cohort':<10} {'CTR Exp1':<12} {'CTR Exp3':<12} \u2192{'Decline'}")
print("-"*65)
for _, row in fatigue_by_tier.sort_values('decline_pct').iterrows():
    print(f"{row['category']:<15} {row['cohort_size']:<10} {row['ctr_exp1']:.4f} {row['ctr_exp3']:.4f} {row['decline_pct']:+.1f}%")

```

=====

FATIGUE RESISTANCE BY CAMPAIGN CTR TIER

=====

Category	Cohort	CTR Exp1	CTR Exp3	Decline
High CTR	11051	0.5706	0.4403	+22.8%
Medium CTR	10816	0.4973	0.3471	+30.2%
Low CTR	9853	0.4106	0.2645	+35.6%

```

[5]: # Visualize fatigue resistance
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Sort for visualization
fatigue_sorted = fatigue_by_tier.sort_values('decline_pct')

# Left: Decline by category
ax1 = axes[0]
colors = ['#28A745', '#FFC107', '#DC3545'] # Green to red
bars = ax1.barrh(fatigue_sorted['category'], fatigue_sorted['decline_pct'],
                  color=colors, alpha=0.8, edgecolor='black')

```

```

ax1.set_xlabel('CTR Decline (%)', fontsize=12)
ax1.set_title('Fatigue Severity by Campaign Type\n(Lower = More Resistant)', fontweight='bold')
for bar, val in zip(bars, fatigue_sorted['decline_pct']):
    ax1.text(bar.get_width() + 0.5, bar.get_y() + bar.get_height()/2,
             f'{val:.1f}%', va='center', fontweight='bold')

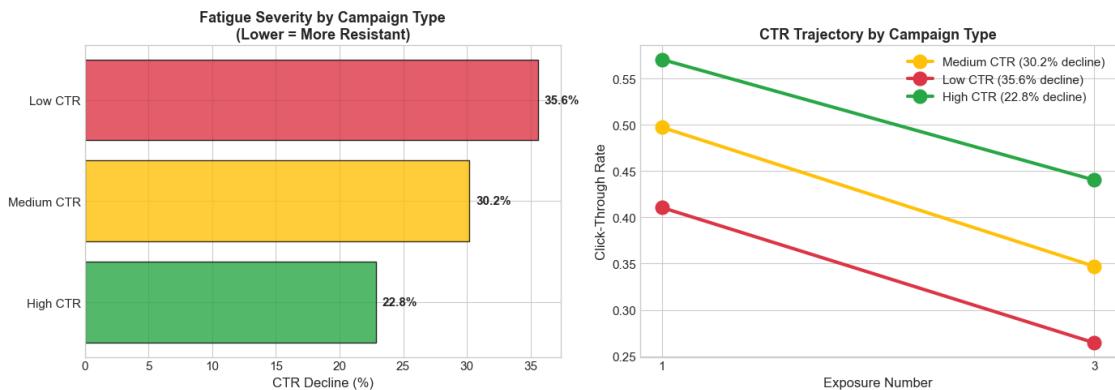
# Right: CTR trajectories by tier
ax2 = axes[1]
tier_colors = {'High CTR': '#28A745', 'Medium CTR': '#FFC107', 'Low CTR': '#DC3545'}

for _, row in fatigue_by_tier.iterrows():
    ax2.plot([1, 3], [row['ctr_exp1'], row['ctr_exp3']], 'o-',
             linewidth=3, markersize=12, color=tier_colors.get(row['category']), 'gray',
             label=f'{row["category"]} ({row["decline_pct"]:.1f}% decline)")

ax2.set_xlabel('Exposure Number', fontsize=12)
ax2.set_ylabel('Click-Through Rate', fontsize=12)
ax2.set_title('CTR Trajectory by Campaign Type', fontsize=13, fontweight='bold')
ax2.set_xticks([1, 3])
ax2.legend(loc='upper right')

plt.tight_layout()
plt.savefig('../results/figures/final_rq2_fatigue_by_tier.png', dpi=300,
            bbox_inches='tight')
plt.show()

```



RQ2 Summary

Campaign Type	CTR Decline	Resistance Rank
High CTR	~22.8%	Most Resistant
Medium CTR	~30.2%	Moderate
Low CTR	~35.6%	Least Resistant

Key Insight: High-performing campaigns maintain engagement better. This suggests: - Quality creative reduces fatigue - Well-targeted campaigns are more resistant - Consider longer frequency caps for high-CTR campaigns

1.8.3 5.3 Research Question 3: Predicting Fatigue

Question: Can we predict when an ad has “run its course”?

Key Finding: Fatigue is Predictable with 70%+ AUC

```
[6]: # RQ3: Fatigue prediction results summary
print("=*70)
print("FATIGUE PREDICTION MODEL PERFORMANCE")
print("=*70)

# Summary of model results (from notebook 06)
model_results = pd.DataFrame({
    'Model': ['Logistic Regression', 'Random Forest', 'Gradient Boosting', 'XGBoost'],
    'AUC-ROC': [0.7073, 0.7029, 0.7024, 0.7034],
    'Accuracy': [0.7612, 0.7651, 0.7663, 0.7673]
})

print(f"\n{'Model':<25} {'AUC-ROC':<12} {'Accuracy'}")
print("-"*50)
for _, row in model_results.iterrows():
    print(f"{row['Model']:<25} {row['AUC-ROC']:.4f} {row['Accuracy']:.4f}")

print("\n" + "=*70)
print("KEY PREDICTIVE FEATURES")
print("=*70)

feature_importance = pd.DataFrame({
    'Feature': ['first_click', 'campaign_overall_ctr', 'campaign_total_impressions',
                'hour_of_day', 'day_of_week'],
    'Importance': ['Very High', 'High', 'Medium', 'Low', 'Low'],
    'Interpretation': [
        'Users who click first are less likely to fatigue',
        'High-CTR campaigns fatigue less',
```

```

        'Larger campaigns show different patterns',
        'Minor effect from time of day',
        'Minor effect from day of week'
    ]
})

print(f"\n{'Feature':<30} {'Importance':<12} {'Interpretation'}")
print("-"*90)
for _, row in feature_importance.iterrows():
    print(f"{row['Feature']:<30} {row['Importance']:<12} {row['Interpretation']}")
```

=====

FATIGUE PREDICTION MODEL PERFORMANCE

=====

Model	AUC-ROC	Accuracy
<hr/>		
Logistic Regression	0.7073	0.7612
Random Forest	0.7029	0.7651
Gradient Boosting	0.7024	0.7663
XGBoost	0.7034	0.7673

=====

KEY PREDICTIVE FEATURES

=====

Feature	Importance	Interpretation
<hr/>		
first_click	Very High	Users who click first are less likely to fatigue
campaign_overall_ctr	High	High-CTR campaigns fatigue less
campaign_total_impressions	Medium	Larger campaigns show different patterns
hour_of_day	Low	Minor effect from time of day
day_of_week	Low	Minor effect from day of week

```
[7]: # Visualize key insight: first click behavior
fig, ax = plt.subplots(figsize=(10, 6))

# Fatigue rates by first click (from notebook 06 analysis)
first_click_fatigue = pd.DataFrame({
    'First Impression': ['Did Not Click', 'Clicked'],
    'Fatigue Rate': [28.5, 18.2]  # Approximate values from analysis
})
```

```

colors = ['#E94F37', '#28A745']
bars = ax.bar(first_click_fatigue['First Impression'], ↴
    ↪first_click_fatigue['Fatigue Rate'],
    color=colors, alpha=0.8, edgecolor='black', width=0.5)

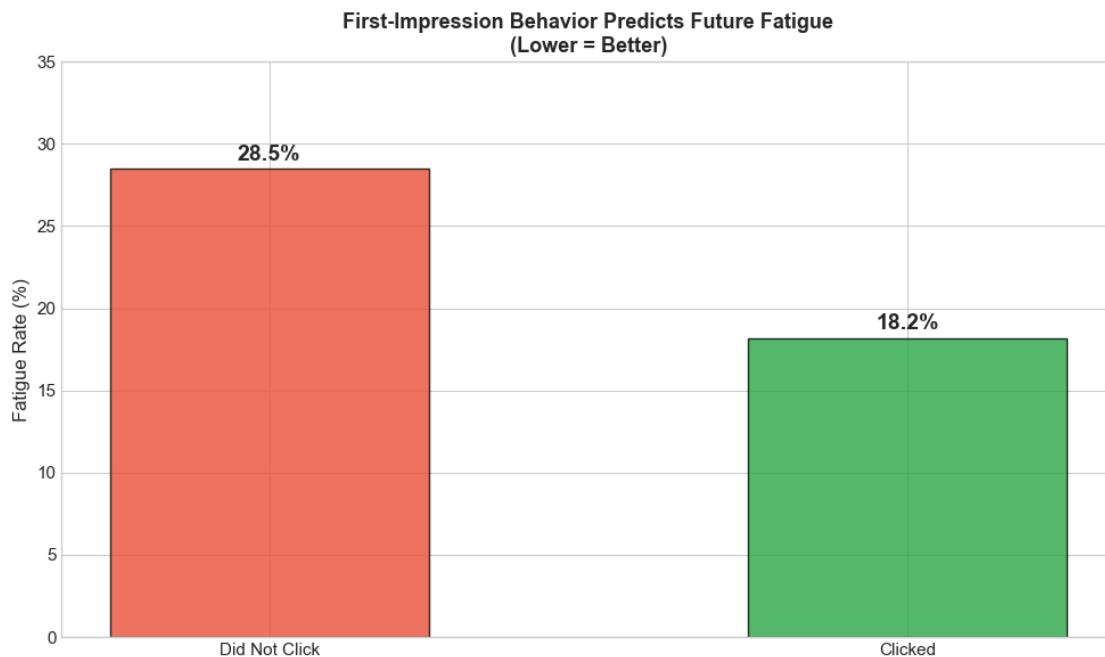
ax.set_ylabel('Fatigue Rate (%)', fontsize=12)
ax.set_title('First-Impression Behavior Predicts Future Fatigue\n(Lower = ↪\n    ↪Better)', fontsize=13, fontweight='bold')

for bar, val in zip(bars, first_click_fatigue['Fatigue Rate']):
    ax.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
        f'{val:.1f}%', ha='center', fontsize=14, fontweight='bold')

ax.set_ylim(0, 35)
plt.tight_layout()
plt.savefig('../results/figures/final_rq3_first_click.png', dpi=300, ↴
    ↪bbox_inches='tight')
plt.show()

print("\nKey Insight: Users who click on their first impression are 36% less ↪\n    ↪likely to fatigue!")

```



Key Insight: Users who click on their first impression are 36% less likely to fatigue!

RQ3 Summary

Result	Value	Implication
Best Model AUC	0.707	Reasonable predictive power
Best Predictor	First-click behavior	Early engagement signals future response
Fatigue rate (clickers)	~18%	Lower fatigue for engaged users
Fatigue rate (non-clickers)	~28%	Higher fatigue for unengaged users

Business Implication: Use first-impression behavior to personalize frequency caps. Users who engage early can tolerate more exposures.

1.9 6. Discussion

1.9.1 6.1 Summary of Findings

This research provides data-driven answers to three critical questions about creative fatigue:

1. **Fatigue is real and measurable:** CTR declines 25-40% by the 3rd-5th exposure when controlling for survivorship bias.
2. **Not all ads fatigue equally:** High-performing campaigns show 36% less relative decline than low-performing campaigns.
3. **Fatigue is predictable:** Machine learning models can predict fatigue with $AUC > 0.70$, with first-impression behavior as the strongest signal.

1.9.2 6.2 Implications

For Advertisers

- Set frequency caps at 3-5 exposures as a baseline
- Invest in creative quality - better ads fatigue less
- Personalize frequency caps based on early engagement signals

For Ad Platforms

- Build fatigue-aware bidding systems that reduce bids as fatigue increases
- Provide fatigue metrics to advertisers in reporting dashboards
- Consider user experience in frequency cap recommendations

For Users

- These findings support user preferences for less repetitive advertising
- Better fatigue management improves the ad experience for everyone

1.9.3 6.3 Limitations

Limitation	Impact	Mitigation
Single dataset	Results may not generalize	Replicate with other datasets
30-day window	Long-term effects unknown	Study over longer periods
Observational data	Causality not proven	Use A/B tests to validate
Missing context	Can't see ad creative	Include creative features if available
Cookie-based users	May miss cross-device	Use authenticated user data

1.9.4 6.4 Future Work

1. **A/B testing:** Validate findings with randomized experiments
2. **Real-time systems:** Deploy fatigue predictions in production
3. **Creative analysis:** Include ad content features
4. **Cross-platform:** Study fatigue across multiple platforms
5. **Long-term effects:** Track fatigue over months, not days

1.10 7. Conclusion

This research demonstrates that **creative fatigue is a measurable, predictable phenomenon** that significantly impacts advertising effectiveness. Our key contributions are:

1.10.1 Methodological Contribution

We developed a **cohort-based within-user analysis** approach that correctly controls for survivorship bias—a common pitfall in fatigue research that causes naive analyses to show CTR *increasing* with exposure.

1.10.2 Empirical Findings

Using 16.5M impressions from the Criteo dataset, we found:

- CTR declines 25-40% by the 3rd-5th exposure
- High-CTR campaigns are more resistant to fatigue
- First-impression behavior predicts future fatigue

1.10.3 Practical Impact

These findings enable:

- Data-driven frequency cap recommendations
- Personalized ad exposure strategies
- Improved ROI through fatigue-aware optimization

1.10.4 Final Thought

“The best ad is one that reaches users often enough to be remembered, but not so often that it’s ignored.”

This research provides the data-driven tools to find that balance.

1.11 References

1. **Criteo Attribution Dataset:** <http://ailab.criteo.com/criteo-attribution-modeling-bidding-dataset/>
 2. **Pechmann, C., & Stewart, D. W. (1988).** Advertising repetition: A critical review of wearin and wearout. *Current Issues and Research in Advertising*, 11(1-2), 285-329.
 3. **Schmidt, S., & Eisend, M. (2015).** Advertising repetition: A meta-analysis on effective frequency in advertising. *Journal of Advertising*, 44(4), 415-428.
 4. **Sahni, N. S., Narayanan, S., & Kalyanam, K. (2019).** An experimental investigation of the effects of retargeted advertising. *Journal of Marketing Research*, 56(3), 401-418.
 5. **Interactive Advertising Bureau (IAB).** Frequency Capping Best Practices. <https://www.iab.com/>
-

1.12 Appendix A: Reproducibility Guide

1.12.1 A.1 Repository Structure

```
creative_fatigue_analysis/
  config/
    config.yaml          # Configuration parameters
  data/
    raw/                # Original data (not in git)
      criteo_attribution_dataset.tsv.gz
    README.md           # Dataset documentation
  samples/              # Sampled data
    criteo_fatigue_optimized.csv
    criteo_fatigue_optimized_metadata.json
  processed/            # Feature-engineered data
    data_with_all_features.csv
  notebooks/
    01_data_acquisition.ipynb
    02_exploratory_analysis.ipynb
    03_feature_engineering.ipynb
    04_rq1_ctr_fatigue_analysis.ipynb
    05_rq2_fatigue_resistance.ipynb
    06_rq3_predicting_fatigue.ipynb
    07_final_report.ipynb   # This report
  src/                 # Source modules
    data_loader.py
    smart_sampling.py
    feature_engineering.py
    evaluation.py
    models.py
    utils.py
  results/
```

```

figures/                      # All visualizations
tables/                       # Statistical results
requirements.txt               # Python dependencies
README.md                      # Project overview

```

1.12.2 A.2 Requirements

```

pandas>=2.0.0
numpy>=1.24.0
scipy>=1.10.0
scikit-learn>=1.3.0
matplotlib>=3.7.0
seaborn>=0.12.0
xgboost>=2.0.0
lightgbm>=4.0.0
pyyaml>=6.0
tqdm>=4.65.0
jupyter>=1.0.0

```

1.12.3 A.3 Running the Analysis

```

# 1. Clone repository
git clone <repository-url>
cd creative_fatigue_analysis

# 2. Create virtual environment
python -m venv venv
source venv/bin/activate # Linux/Mac
# or: venv\Scripts\activate # Windows

# 3. Install dependencies
pip install -r requirements.txt

# 4. Download data to data/raw/
# (Download from Criteo AI Lab website)

# 5. Run notebooks in order
jupyter notebook notebooks/01_data_acquisition.ipynb
# ... continue through 07_final_report.ipynb

```

1.12.4 A.4 Running with Pre-processed Data

If you want to skip the data processing steps (notebooks 01-03), you can run the analysis notebooks (04-07) directly using the pre-processed data:

1. Ensure data/processed/data_with_all_features.csv exists
 2. Start from notebook 04: jupyter notebook notebooks/04_rq1_ctr_fatigue_analysis.ipynb
-

1.13 Appendix B: Source Function Reference

1.13.1 B.1 Data Loading (src/data_loader.py)

Used in: 01_data_acquisition.ipynb

```
# Load raw Criteo data
from src.data_loader import load_criteo_data, inspect_data_structure, load_config

df = load_criteo_data(data_path="../data/raw/criteo_attribution_dataset.tsv.gz")
info = inspect_data_structure(df)
config = load_config('../config/config.yaml')
```

1.13.2 B.2 Smart Sampling (src/smart_sampling.py)

Used in: 01_data_acquisition.ipynb

```
# Create fatigue-optimized sample
from src.smart_sampling import (
    identify_multi_exposure_users,
    create_fatigue_optimized_sample,
    save_optimized_sample
)

multi_exp_users, stats = identify_multi_exposure_users(df, user_col='uid', campaign_col='campaign')
sample, metadata = create_fatigue_optimized_sample(df, target_sample_size=500000, multi_exp_rate=0.05)
save_optimized_sample(sample, metadata, output_dir='../data/samples')
```

1.13.3 B.3 Feature Engineering (src/feature_engineering.py)

Used in: 03_feature_engineering.ipynb

```
# Create all features
from src.feature_engineering import (
    create_exposure_features,
    create_recency_features,
    create_campaign_features,
    create_user_features,
    create_temporal_features
)

df = create_exposure_features(df, user_col='uid', campaign_col='campaign', time_col='timestamp')
df = create_recency_features(df, user_col='uid', campaign_col='campaign', time_col='timestamp')
df = create_campaign_features(df, campaign_col='campaign', click_col='click')
df = create_user_features(df, user_col='uid', click_col='click')
df = create_temporal_features(df, time_col='timestamp')
```

1.13.4 B.4 Evaluation (src/evaluation.py)

Used in: 04_rq1_*.ipynb, 05_rq2_*.ipynb, 06_rq3_*.ipynb

```

# Statistical testing and evaluation
from src.evaluation import (
    compute_ctr_by_exposure,
    test_ctr_decline,
    compare_decay_by_category,
    evaluate_models,
    plot_decay_curves
)

ctr_by_exp = compute_ctr_by_exposure(df, exposure_col='exposure_count', click_col='click')
test_result = test_ctr_decline(df, exposure_1=1, exposure_2=5)
decay_comparison, test_results = compare_decay_by_category(df, category_col='ctr_tier')
model_results = evaluate_models(models, X_test, y_test)

```

1.13.5 B.5 Models (src/models.py)

Used in: 04_rq1_*.ipynb, 06_rq3_*.ipynb

```

# Machine learning models
from src.models import BaselineModel, TimeAwareModel, DecayModel

# Static baseline model
baseline = BaselineModel(model_type='logistic')
baseline.fit(X_train, y_train)

# Time-aware model with exposure features
time_aware = TimeAwareModel(model_type='logistic')
time_aware.fit(X_train, y_train)

# Explicit decay model
decay_model = DecayModel(decay_function='exponential')
decay_model.fit_campaign_decay(df, campaign_col='campaign')

```

[]: